

Summary: Quantitative Risk Management

Fabian MARBACH, Spring Semester 2016

1 Risk in perspective

Definition of risk

- hazard, a chance of bad consequences, loss or exposure to mischance
- any event or action that may adversely affect an organization's ability to achieve its objectives and execute its strategies

Financial risks Good risk management has to follow a holistic approach, i.e. all of the following types of risks and their interactions should be considered.

- **Market risk:** risk of loss in a financial position due to changes in the *underlying* components, e.g. stocks, bonds, commodity prices
- **Credit risk:** risk of a *counterparty* failing to meet its obligations (*default*), i.e. the risk of not receiving promised repayments, e.g. loans, bonds
- **Operational risk:** risk of loss resulting from inadequate or *failed internal processes, people and systems* or from *external events*, e.g. fraud, earthquakes
- **Liquidity risk:**
 - *Market liquidity risk:* lack of marketability of an investment that cannot be bought or sold quickly enough to prevent/minimize a loss
 - *Funding liquidity risk:* refers to the ease with which institutions can raise funding
- **Underwriting risk:** in insurance, the risk inherent to *insurance policies sold*, e.g. natural catastrophes, political changes, etc.
- **Model risk:** using a *misspecified (inappropriate) model* for measuring risk

Risk measurement Assume a portfolio of d investments with weights w_1, \dots, w_d . Denote by X_j the change in value of the j^{th} investment. Then the change in value (profit and loss, P&L) of the portfolio is:

$$X = \sum_{j=1}^d w_j X_j$$

Measuring risk consists then of determining the distribution function F for the joint model of $\mathbf{X} = (X_1, \dots, X_d)$.

Interpretation: A risk measure can be interpreted as the amount of capital that needs to be added to a position so that it becomes acceptable to the regulator.

Risk management is a discipline for living with the possibility that future events may cause adverse effects. It is about ensuring *resilience to future events*. It involves:

- Determine *capital to hold to absorb losses* (due to regulatory and economic capital purposes).
- Ensure portfolios to be *well diversified*.
- *Optimize portfolios* according to risk-return considerations.

Three-pillar concept

- **Minimal capital charge:** calculate minimum regulatory capital to ensure that a bank holds *sufficient capital* for its *market risk* in the trading book, *credit risk* in the banking book and *operational risk*.
- **Supervisory review process:** local regulators conduct capital adequacy assessments (reviews, stress tests).
- **Market discipline:** better public disclosure of risk measures and other relevant information.

2 Basic concepts in risk management

2.1 Modelling value and value change

Value and loss

- V_t : value of a portfolio of assets and possibly liabilities
- Δt : time horizon
 - portfolio composition remains fixed during Δt
 - no intermediate payments during Δt \leadsto fine for small Δt but unlikely to hold for large Δt
- $\mathbf{Z} = (Z_{t,1}, \dots, Z_{t,d}) \in \mathbb{R}^d$: **risk factors**
- $\mathbf{X}_{t+1} = \mathbf{Z}_{t+1} - \mathbf{Z}_t$: **risk-factor changes**
- **value** (mapping of risks) and change in value:

$$V_t = f(t, \mathbf{Z}_t) \quad \Delta V_{t+1} = V_{t+1} - V_t$$

For longer time intervals: $\Delta V_{t+1} = V_{t+1}/(1+r) - V_t$ with r the risk-free interest rate.

■ Loss:

$$\begin{aligned} L_{t+1} &= -\Delta V_{t+1} \\ &= -(f(t+1, \mathbf{Z}_{t+1}) - f(t, \mathbf{Z}_t)) \\ &= -(f(t+1, \mathbf{Z}_t + \mathbf{X}_{t+1}) - f(t, \mathbf{Z}_t)) \end{aligned}$$

The distribution of L_{t+1} is called *loss distribution* and is determined by the loss distribution of risk-factor changes \mathbf{X}_{t+1} . The *profit-and-loss (P&L) distribution* is the distribution of $-L_{t+1} = \Delta V_{t+1}$.

■ Linearized loss:

$$\begin{aligned} L_{t+1}^\Delta &= - \left(\underbrace{f_t(t, \mathbf{Z}_t)}_{=:c_t} + \sum_{j=1}^d \underbrace{f_{z_j}(t, \mathbf{Z}_t)}_{=:b_{t,j}} \cdot X_{t+1,j} \right) \\ &= -(c_t + \mathbf{b}_t' \mathbf{X}_{t+1}) \end{aligned}$$

The approximation is best if the *risk-factor changes are small in absolute value*.

Remark: The Taylor approximation presumes that no sudden large movements occur in the risk-factor changes \mathbf{X}_{t+1} within a short time-period $\Delta t = 1$, which of course does not always hold in reality.

Examples

■ Stock portfolio

- portfolio of stocks $S_{t,1}, \dots, S_{t,d}$ with λ_j the number of shares in stock j
- risk factors: log-prices $Z_{t,j} = \log S_{t,j}$
- value: $V_t = \sum_{j=1}^d \lambda_j e^{Z_{t,j}}$
- one-period ahead loss: $L_{t+1} = - \sum_{j=1}^d \lambda_j \underbrace{S_{t,j}}_{=:w_{t,j}} (e^{X_{t+1,j}} - 1)$
- linearized loss: $L_{t+1}^\Delta = -\mathbf{w}_t' \mathbf{X}_{t+1}$
- with $\boldsymbol{\mu} = \mathbb{E}[\mathbf{X}_{t+1}]$ and $\boldsymbol{\Sigma} = \text{Cov}[\mathbf{X}_{t+1}]$, the *expectation* and *variance* of the one-period ahead loss are given by:

$$\mathbb{E}[L_{t+1}^\Delta] = \mathbf{w}_t' \boldsymbol{\mu} \quad \text{Var}[L_{t+1}^\Delta] = \mathbf{w}_t' \boldsymbol{\Sigma} \mathbf{w}_t$$

■ European call option

- risk factors: $\mathbf{Z}_t = (\log S_t, r_t, \sigma_t)$

- risk-factor changes:

$$\mathbf{X}_{t+1} = (\log(S_{t+1}/S_t), r_{t+1} - r_t, \sigma_{t+1} - \sigma_t)$$

- value:

$$\begin{aligned} V_t &= C^{\text{BS}}(t, S_t, r, \sigma, K, T) \\ &= C^{\text{BS}}(t, e^{Z_{t,1}}, Z_{t,2}, Z_{t,3}, K, T) = f(t, \mathbf{Z}_t) \end{aligned}$$

- linearized loss:

$$L_{t+1}^{\Delta} = -(C_t^{\text{BS}} \Delta t + C_{S_t}^{\text{BS}} S_t X_{t+1,1} + C_{r_t}^{\text{BS}} X_{t+1,2} + C_{\sigma_t}^{\text{BS}} X_{t+1,3})$$

Note that the *Greeks* enter here.

- For portfolios of derivatives, L_{t+1}^{Δ} can be a rather poor approximation to L_{t+1} , and thus higher-order (second-order Taylor) approximations might be needed (e.g. delta-gamma approximation).

Fair value accounting

Level 1 Mark-to-market: use *quoted prices* for the *same* instrument

Level 2 Mark-to-model w/ objective inputs: use *quoted prices* for *similar* instruments or use valuation techniques/models with inputs based on *observable market data*

Level 3 Mark-to-model w/ subjective inputs: use valuation techniques/models for which some inputs are *not observable* in the market (e.g. loans to companies for which no CDS spreads are available)

Risk-neutral valuation

- value of a financial instrument today = expected discounted values of future cash flows w.r.t. the *risk-neutral pricing measure* \mathbb{Q} (as opposed to the *real world/historical measure* \mathbb{P})

- **risk-neutral pricing rule:**

$$V_t^H = \mathbb{E}_{\mathbb{Q}}[e^{-r(T-t)} H | \mathcal{F}_t], \quad t < T$$

- \mathbb{Q} is calibrated to market prices, while \mathbb{P} is estimated from historical data.

Key statistical tasks of QRM

- Find a *statistical model* for \mathbf{X}_{t+1} (i.e. a model for forecasting \mathbf{X}_{t+1} based on historical data).
- Compute/derive the *PDF/CDF* $F_{L_{t+1}}$ (requires the PDF/CDF of $f(t+1, \mathbf{Z}_t + \mathbf{X}_{t+1})$).
- Compute a *risk measure* from $F_{L_{t+1}}$.

Methods

■ Analytical method (\leadsto variance-covariance method)

- *Idea:* choose $F_{\mathbf{X}_{t+1}}$ and f s.t. $F_{L_{t+1}}$ can be determined explicitly.
prime example: variance-covariance method
- *Assumption 1:* $\mathbf{X}_{t+1} \sim \mathcal{N}(\boldsymbol{\mu}, \boldsymbol{\Sigma})$
(e.g. if \mathbf{Z}_t is a Brownian motion, S_t is a geometric Brownian motion)
- *Assumption 2:* $F_{L_{t+1}^{\Delta}}$ is a good approximation to $F_{L_{t+1}}$
implies: $L_{t+1}^{\Delta} = -(c_t + \mathbf{b}_t' \mathbf{X}_{t+1})$
 $\Rightarrow L_{t+1}^{\Delta} \sim \mathcal{N}(-c_t - \mathbf{b}_t' \boldsymbol{\mu}, \mathbf{b}_t' \boldsymbol{\Sigma} \mathbf{b}_t)$
- *Advantages/Drawbacks:*
 - + $F_{L_{t+1}^{\Delta}}$ explicit (typically risk measures)
 - + easy to implement
 - *linear loss operator* might be a bad approximation
 - assumption of *i.i.d.* distribution might not hold
 - assumption of a *multivariate normal* distribution might be too crude (since it underestimates the tail of $F_{L_{t+1}}$)

■ Historical simulation

- *Idea:* estimate $F_{L_{t+1}}$ by its *empirical distribution function* (EDF), i.e.

$$\tilde{F}_{L_{t+1}}(x) = \frac{1}{n} \sum_{i=1}^n \mathbb{I}_{\{\tilde{L}_{t-i+1} \leq x\}}$$

based on $\tilde{L}_k = -(f(t+1, \mathbf{Z}_t + \mathbf{X}_k) - f(t, \mathbf{Z}_t))$.
 $\tilde{L}_{t-n+1}, \dots, \tilde{L}_t$ show what would happen to the current portfolio if the past n risk-factor changes were to recur.

- *Advantages/Drawbacks:*
 - + does not require a *joint model* of the risk factors \mathbf{X}_{t+1}
 \leadsto no estimation of the distribution of \mathbf{X}_{t+1} required (estimating the dependencies is usually the most challenging)
 - + easy to implement
 - sufficient relevant and synchronized *data* for all risk-factor changes required \leadsto sample size might be too small
 - historical data may not contain (sufficient) examples of *extreme scenarios*
 - considers only past losses (\leadsto "driving a car by looking in the back mirror")

■ Monte Carlo method

- *Idea:* take any model for \mathbf{X}_{t+1} , simulate from it, compute the corresponding simulated losses and estimate $F_{L_{t+1}}$ (typically the EDF)
- *Advantages/Drawbacks:*

- + sample size and number of repetitions can be increased freely (\Rightarrow risk measures can be estimated with greater accuracy)
- + quite general, i.e. applicable to any model of \mathbf{X}_{t+1} which is easy to sample
- does not solve the *problem of finding a joint distribution* of the risk factors (i.e. it is still unclear how to find an appropriate model for \mathbf{X}_{t+1}) \leadsto any result is only as good as the chosen $F_{\mathbf{X}_{t+1}}$
- *computational cost*, i.e. every simulation requires to evaluate the portfolio, (e.g. *Nested Monte Carlo simulations*: especially expensive if the portfolio contains derivatives which are priced via Monte Carlo themselves)

2.2 Risk measurement

Approaches to risk measurement

■ Notional-amount approach

- *risk of a portfolio* = summed notional values of the securities \times their riskiness factor
- *Advantages:* simplicity
- *Drawbacks:* no differentiation between long and short positions, no netting (eg. risk of a hedged position is twice the risk of an unhedged position)

■ Risk measures based on loss distributions

- *risk of a portfolio:* a characteristic of the underlying *loss distribution* over some time horizon Δt
- e.g. variance, VaR, ES
- *Advantages:* this concept makes sense on all levels, e.g. reflects netting and diversification effects (if estimated properly)
- *Drawbacks:* estimates of loss distributions based on *past data* (backward-looking), loss distributions are difficult to estimate \leadsto could be completed by information from *scenarios* (forward-looking)

■ Scenario-based risk measures

- *risk of a portfolio:* maximum weighted loss under all relevant scenarios
If $\chi = \{\mathbf{x}_1, \dots, \mathbf{x}_n\}$ denote the risk-factor changes (scenarios) with corresponding weights $\mathbf{w} = (w_1, \dots, w_n)$, the risk is:

$$\begin{aligned} \psi_{\chi, \mathbf{w}} &= \max_{1 \leq i \leq n} \{w_i L(\mathbf{x}_i)\} \\ &= \max_{1 \leq i \leq n} \{\mathbb{E}_{\mathbb{P}_i}[L(\mathbf{X})] : \mathbf{X} \sim \mathbb{P} \in \{\mathbb{P}_1, \dots, \mathbb{P}_n\}\} \end{aligned}$$

where $L(x)$ denotes the loss the portfolio would suffer if the hypothetical scenario x were to occur and where $\mathbf{X}_i \sim \mathbb{P}_i = w_i \delta_{\mathbf{x}_i} + (1 - w_i) \delta_{\mathbf{0}}$ is a probability measure on \mathbb{R}^d .

Such a risk measure is known as a *generalized scenario*.

- **Advantages:** useful for portfolios with *few risk factors*, use full *complementary information* to risk measures based on loss distributions (past data)
- **Drawbacks:** determining scenarios and weights

2.2.1 Risk measures

Coherent risk measure Assume $L, L_1, L_2 \in \mathcal{M}$, where \mathcal{M} a linear space of random variables.

Axiom 1 Monotonicity

$$L_1 \leq L_2 \Rightarrow \rho(L_1) \leq \rho(L_2)$$

- **Interpretation:** positions which lead to a higher loss in every state of the world require more risk capital.

Axiom 2 Translation invariance

$$\rho(L + l) = \rho(L) + l, \quad \forall l \in \mathbb{R}$$

- **Interpretation:** by adding a (cash position) l to a position with loss L , we alter the capital requirements accordingly.

- **Criticism:** most people believe this to be reasonable.

Axiom 3 Subadditivity

$$\rho(L_1 + L_2) \leq \rho(L_1) + \rho(L_2)$$

- **Interpretation:** reflects the idea of *diversification*, using a non-subadditive ρ encourages institutions to legally break up into subsidiaries to reduce regulatory capital requirements, subadditivity makes decentralization possible.

- **Criticism:** VaR is ruled out under certain scenarios.

Axiom 4 Positive homogeneity

$$\rho(\lambda L) = \lambda \rho(L), \quad \forall \lambda > 0$$

- **Interpretation:** n times the same loss means no diversification, so equality should hold.

- **Criticism:** if λ is large, *liquidity risk* plays a role and one should rather have $\rho(\lambda L) > \lambda \rho(L)$ (also to penalize concentration of risk), but this contradicts subadditivity (\leadsto convex risk measures).

One can show that all coherent risk measures can be represented as **generalized scenarios** via:

$$\rho(L) = \sup\{\mathbb{E}_{\mathbb{P}}[L] : \mathbb{P} \in \mathcal{P}\}$$

for a suitable set \mathcal{P} of probability measures.

Convex risk measures A risk measure ρ that is *monotone*, *translation invariant* and *convex* is called a convex risk measure.

- Any coherent risk measure is also a convex risk measure, while the converse is in general not true.

Value-at-risk (VaR)

- For a loss $L \sim F_L$, **value-at-risk (VaR)** at *confidence level* $\alpha \in (0, 1)$ is defined by:

$$\text{VaR}_{\alpha}(L) = F_L^{\leftarrow}(\alpha) = \inf\{x \in \mathbb{R} : F_L(x) \geq \alpha\}$$

- VaR_{α} is the α -quantile of F_L . It thus holds:

$$F_L(x) < \alpha \quad \forall x < \text{VaR}_{\alpha}(L)$$

$$F_L(\text{VaR}_{\alpha}(L)) = F_L(F_L^{\leftarrow}(\alpha)) \geq \alpha$$

- VaR_{α} under some distributions:

- Exponential distribution ($X \sim \text{Exp}(\lambda)$, $F(x) = 1 - e^{-\lambda x}$):
 $\text{VaR}_{\alpha}(X) = -\frac{1}{\lambda} \log(1 - \alpha)$

- Pareto distribution ($X \sim \text{Pareto}(\lambda)$, $F(x) = 1 - x^{-\lambda}$):
 $\text{VaR}_{\alpha}(X) = (1 - \alpha)^{-\frac{1}{\lambda}}$

- VaR_{α} as a **coherent risk measure**:

- VaR is **in general not a coherent risk measure** since it is *in general not subadditive* (but it fulfills the other three requirements of a coherent risk measure).

Examples:

- (i) If X_i are *highly skewed*

E.g. let X_i for $i = 1, \dots, 100$ be i.i.d. RVs s.t.

$$X_i = \begin{cases} -2 & \text{with probability 0.99} \\ 100 & \text{with probability 0.01} \end{cases}$$

- (ii) If X_i have *infinite mean*

E.g. let X_1, X_2 be independent RVs with $\mathbb{P}[X_i \leq x] = 1 - x^{-1/2}$, $\forall x \geq 1$, for $i = 1, 2$.

- $\text{VaR}_{\alpha}(X, Y)$ is subadditive and thus a *coherent risk measure* in the following cases:

- (i) if X, Y are *comonotonic* then $\text{VaR}_{\alpha}(X, Y)$ is *additive*. Thus, VaR_{α} is *comonotone additive*. (i.e. $\text{VaR}_{\alpha}(X + Y) = \text{VaR}_{\alpha}(X) + \text{VaR}_{\alpha}(Y)$)

- (ii) for $\alpha \geq 0.5$: if X, Y are *elliptically distributed* (e.g. normal distribution)

- It holds in general for two RVs X, Y that:

$$\text{VaR}_{\alpha}(X, Y) \leq \text{ES}_{\alpha}(X) + \text{ES}_{\alpha}(Y)$$

- Fallacy w.r.t. VaR and linear correlation:

For two RVs X, Y with finite second moment, $\text{VaR}_{\alpha}(X + Y)$ is *not* maximal if the linear correlation between X, Y (i.e. $\text{Corr}[X, Y]$) is maximal since VaR is in general not a coherent risk measure (i.e. some X, Y can be found s.t. $\text{VaR}_{\alpha}(X + Y) > \text{VaR}_{\alpha}(X) + \text{VaR}_{\alpha}(Y)$).

- **Remarks:**

- VaR is the most widely used risk measure (e.g. by Basel II or Solvency II).
- $\text{VaR}_{\alpha}(L)$ is *not* a *what-if* risk measure, i.e. it does not provide information about the severity of losses which occur with probability $\leq 1 - \alpha$.

Expected shortfall (ES)

- For a loss $L \sim F_L$ with $\mathbb{E}[|L|] < \infty$, **expected shortfall (ES)** at *confidence level* $\alpha \in (0, 1)$ is defined by:

$$\begin{aligned} \text{ES}_{\alpha}(L) &= \frac{1}{1 - \alpha} \int_{\alpha}^1 \text{VaR}_u(L) du \\ &= \mathbb{E}[L | L \geq \text{VaR}_{\alpha}(L)] = \frac{1}{1 - \alpha} \int_{\text{VaR}_{\alpha}(L)}^{\infty} x f_L(x) dx \end{aligned}$$

where f_L denotes the PDF of L (if it exists!).

- ES_{α} is the **average over** $\text{VaR}_u, \forall u \geq \alpha$. If F_L is continuous, ES_{α} is the average loss beyond VaR_{α} . Thus: $\text{ES}_{\alpha} \geq \text{VaR}_{\alpha}$.
- ES_{α} looks further into the tail of F_L . It is a *what-if* risk measure, i.e. VaR_{α} is *frequency-based*, ES_{α} is *severity-based*.
- ES_{α} is a *coherent risk measure* (for continuous RVs!). ES_{α} is *comonotone additive*.

- **In practice:**

Besides VaR, ES is the most important risk measure. ES_{α} is more difficult to estimate and backtest than VaR_{α} since a larger sample size is required and the variance of estimators is typically larger.

Risk measures under the normal distribution \mathcal{N}

- Value-at-Risk VaR_{α}

- For $X \sim \mathcal{N}(\mu, \sigma^2)$: $\text{VaR}_{\alpha}(X) = \mu + \sigma \Phi^{-1}(\alpha)$

- For $X_1, \dots, X_n \stackrel{\text{i.i.d.}}{\sim} \mathcal{N}(\mu, \sigma^2)$, it follows that $X_1 + \dots + X_n \sim \mathcal{N}(n\mu, n\sigma^2)$ and thus:

$$\text{VaR}_{\alpha}(X_1 + \dots + X_n) = n\mu + \sqrt{n}\sigma \Phi^{-1}(\alpha)$$

- For $Y \sim \mathcal{N}(\mathbf{a}^{\top} \boldsymbol{\mu}, \mathbf{a}^{\top} \boldsymbol{\Sigma} \mathbf{a})$:

$$\text{VaR}_{\alpha}(Y) = \mathbf{a}^{\top} \boldsymbol{\mu} + \sqrt{\mathbf{a}^{\top} \boldsymbol{\Sigma} \mathbf{a}} \Phi^{-1}(\alpha)$$

- Expected Shortfall ES_{α}

- For $X \sim \mathcal{N}_1(\mu, \sigma^2)$:

$$\text{ES}_\alpha(X) = \mu + \sigma \frac{\varphi(\Phi^{-1}(\alpha))}{1 - \alpha}$$

and it holds that:

$$\lim_{\alpha \uparrow 1} \frac{\text{ES}_\alpha(X)}{\text{VaR}_\alpha(X)} = 1, \quad \text{VaR}_\alpha(X) \leq \text{ES}_\alpha(X)$$

- For $Y \sim \mathcal{N}_1(\mathbf{a}^\top \mu, \mathbf{a}^\top \Sigma \mathbf{a}) \Leftrightarrow \mathbf{Y} = \mathbf{a}^\top \mathbf{X}, \mathbf{X} \sim \mathcal{N}_d(\mu, \Sigma)$:

$$\text{ES}_\alpha(Y) = \mathbf{a}^\top \mu + \sqrt{\mathbf{a}^\top \Sigma \mathbf{a}} \frac{\varphi(\Phi^{-1}(\alpha))}{1 - \alpha}$$

3 Empirical properties of financial data

Stylized facts about univariate financial return series

- (U1) Return series are *not i.i.d.* although they show *little serial correlation*.
- (U2) Series of *absolute* or squared returns show *profound serial correlation*.
- (U3) Conditional *expected returns* are close to zero.
- (U4) *Volatility* (conditional standard deviation) appears to *vary over time*.
- (U5) *Extreme returns* appear in *clusters* (volatility clustering).
- (U6) Return series are *leptokurtic* or *heavy-tailed* (power-like tail).

Stylized facts about multivariate financial return series

- (M1) Multivariate return series show little evidence of *cross-correlation*, except for *contemporaneous returns* (i.e. at the same t).
- (M2) Multivariate series of *absolute returns* show *profound cross-correlation*.
- (M3) *Correlations* between contemporaneous returns *vary over time*.
- (M4) *Extreme returns* in one series often *coincide with extreme returns* in several other series (e.g. tail dependence).

4 Financial time series

Remark: not part of the exam.

5 Extreme value theory

5.1 Maxima (GEV)

Convergence of sums

- Let $(X_k)_{k \in \mathbb{N}}$ be i.i.d. with $\mathbb{E}[X_1^2] < \infty$ (mean μ , variance σ^2). Define $S_n = \sum_{k=1}^n X_k$.
- As $n \rightarrow \infty$, $\bar{X}_n \xrightarrow{\text{a.s.}} \mu$ by the Strong Law of Large Numbers, so $\frac{\bar{X}_n - \mu}{\sigma} \xrightarrow{\text{a.s.}} 0$.
- By the central limit theorem (CLT):

$$\frac{\sqrt{n}(\bar{X}_n - \mu)}{\sigma} = \frac{S_n - n\mu}{\sqrt{n}\sigma} \xrightarrow{n \uparrow \infty} \mathcal{N}(0, 1)$$

$$\lim_{n \rightarrow \infty} \mathbb{P} \left[\frac{S_n - b_n}{a_n} \leq x \right] = \Phi(x)$$

where the sequences $a_n = \sqrt{n}\sigma$ and $b_n = n\mu$ give normalization.

Block maxima

- Let $(X_i)_{i \in \mathbb{N}} \sim F$ and F continuous. Then the **block maximum** is given by:

$$M_n = \max\{X_1, \dots, X_n\}$$

- For $n \rightarrow \infty$, $M_n \rightarrow x_F$ a.s. where:

$$x_F := \sup\{x \in \mathbb{R} : F(x) < 1\} = F^{\leftarrow}(1) \leq \infty$$

denotes the *right endpoint* of F .

- **Block-maxima method** to compute estimates $\hat{\mu}, \hat{\sigma}, \hat{\xi}$:

- (i) Divide the sample into m blocks of size n ;
- (ii) compute the maximum M_i for each block;
- (iii) fit the GEV distribution $H_\xi\left(\frac{x-\mu}{\sigma}\right)$ to the sample of block maxima M_1, \dots, M_m (e.g. MLE, method of moments).
- If $X_1, \dots, X_n \stackrel{\text{i.i.d.}}{\sim} F$ and $M_n = \max\{X_1, \dots, X_n\}$ or $M_n \sim H$, the k **n -block return level** is:

$$r_{n,k} = H^{\leftarrow}\left(1 - \frac{1}{k}\right) = (F^n)^{\leftarrow}\left(1 - \frac{1}{k}\right)$$

thus: $\mathbb{P}[M_n > r_{n,k}] = \frac{1}{k}$

- $r_{n,k}$ is the level which is expected to be exceeded in one out of every k blocks of size n .
- $r_{n,k}$ is the $(1 - \frac{1}{k})$ quantile of the distribution of M_n .

- **Parametric estimation:**
Approximate $F^n(x) \approx H_\xi\left(\frac{x-\mu}{\sigma}\right) =: H_{\xi,\mu,\sigma}$ for some $\mu \in \mathbb{R}$, $\sigma > 0$. Then:

$$\hat{r}_{n,k} = H_{\hat{\xi},\hat{\mu},\hat{\sigma}}^{\leftarrow}\left(1 - \frac{1}{k}\right) = \hat{\mu} + \frac{\hat{\sigma}}{\hat{\xi}} \left(\left(-\log\left(1 - \frac{1}{k}\right) \right)^{-\hat{\xi}} - 1 \right)$$

The estimates for $\hat{\xi}, \hat{\mu}, \hat{\sigma}$ can be obtained e.g. using MLE.

- For $M_n \sim H$, the **return-period** of the event $\{M_n > u\}$ is $k_{n,u} = \frac{1}{H(u)}$.
 - $k_{n,u}$ is the number of n -blocks for which we expect to see a single n -block exceeding u .
 - Thus, $k_{n,u}$ solves $r_{n,k_{n,u}} = u$.
 - **Parametric estimation:**
As above, approximate $F^n(x) \approx H_\xi\left(\frac{x-\mu}{\sigma}\right) =: H_{\xi,\mu,\sigma}$.
Then: $\hat{k}_{n,u} = 1/\hat{H}_{\hat{\xi},\hat{\mu},\hat{\sigma}}(u)$.
- **Bias-variance trade-off:**
The mean-squared error (MSE) of the estimator $\hat{r}_{n,k}$ can be split into a bias part and a variance part.
Now, IOT determine a reasonable block size, one has to find a compromise between large blocks (which increase the variance of estimates) and small blocks (which induce bias). The fixed relation $N = m \cdot n$ then explains the trade-off.
Remark: The larger the return period $r_{n,k}$, the larger the uncertainty, and thus the wider the confidence interval.

Maximum domain of attraction (MDA)

- F is in the **maximum domain of attraction (MDA)** of H ($F \in \text{MDA}(H)$) if \exists *normalizing sequences* of real numbers $(a_n) > 0$ and $(b_n) \in \mathbb{R}$ s.t. $\frac{M_n - b_n}{a_n}$ converges in distribution to H , i.e.

$$\lim_{n \rightarrow \infty} \mathbb{P} \left[\frac{M_n - b_n}{a_n} \leq x \right] = \lim_{n \rightarrow \infty} F^n(a_n x + b_n) = H(x)$$

i.e. $F^n(x) \simeq H\left(\frac{x - b_n}{a_n}\right)$

for some *non-degenerate density* H (i.e. not a unit jump) and large n .

- **Remarks:**

- In other words, the properly normalized term $(M_n - b_n)/a_n$ converges in distribution to some RV $Z \sim H$.
- One can show that H is determined up to location/scale, i.e. H specifies a unique type of distribution.
This is guaranteed by the convergence to types theorem.
- All commonly applied continuous F belong to $\text{MDA}(H_\xi)$ for some $\xi \in \mathbb{R}$.

Slowly/regularly varying function

- A positive, Lebesgue-measurable function L on $(0, \infty)$ is **slowly varying** at ∞ if:

$$\lim_{x \rightarrow \infty} \frac{L(tx)}{L(x)} = 1, \quad t > 0$$

The class of all such functions is denoted by \mathcal{R}_0 .

– *Examples:* $c, \log \in \mathcal{R}_0$

- A positive, Lebesgue-measurable function h on $(0, \infty)$ is **regularly varying** at ∞ with index $\alpha \in \mathbb{R}$ if:

$$\lim_{x \rightarrow \infty} \frac{h(tx)}{h(x)} = t^\alpha, \quad t > 0$$

The class of all such functions is denoted by \mathcal{R}_α .

– *Examples:* $x^\alpha L(x) \in \mathcal{R}_\alpha$

– *Remark:* If $\bar{F} \in \mathcal{R}_{-\alpha}$, $\alpha > 0$, the tail of F decays like a power function (Pareto like).

Generalized extreme value (GEV) distribution

- The (standard) **generalized extreme value (GEV) distribution** (CDF) is given by:

$$H_\xi(x) = \begin{cases} \exp(-(1 + \xi x)^{-1/\xi}) & : \xi \neq 0 \\ \exp(-e^{-x}) & : \xi = 0 \end{cases}$$

where $1 + \xi x > 0$ (MLE).

- A **three-parameter family** is obtained by the following *location-scale* transform:

$$H_{\xi, \mu, \sigma}(x) = H_\xi\left(\frac{x - \mu}{\sigma}\right), \quad \mu \in \mathbb{R}, \sigma > 0$$

- **Shape parameter:** ξ

Tail index: $\alpha = \frac{1}{\xi}$

The smaller α (the larger ξ), the more heavy-tailed H_ξ and vice-versa.

- A **fitted GEV model** can be used to estimate the:

- size of an event with prescribed frequency (*return-level problem*)
- frequency of an event with prescribed size (*return-period problem*)

- *Remark:*

- The parameterization is continuous in ξ (simplifies statistical modelling).

MDAs for different GEV distribution cases

- $\xi > 0$: **Fréchet MDA** (\leadsto heavy-tailed)

- (i) For $\xi > 0$:

$$F \in \text{MDA}(H_\xi) \iff \bar{F}(x) = x^{-\frac{1}{\xi}} L(x)$$

for some $L \in \mathcal{R}_0$ (i.e. L a slowly varying function at ∞). Thus, $\bar{F}(x)$ has to be regularly varying at ∞ .

- (ii) *Fréchet CDF:*

$$\Phi_\alpha(x) = \begin{cases} \exp(-x^{-\alpha}) & : x > 0 \\ 0 & : x \leq 0 \end{cases} \quad \alpha > 0$$

with shape parameter $\xi = \frac{1}{\alpha}$.

- (iii) *Normalizing sequences:* $a_n = F^{\leftarrow}(1 - \frac{1}{n})$ and $b_n = 0$.

- (iv) *Right endpoint of F :* $x_{H_\xi} = \infty$

- (v) *Moments:* If $X \sim F \in \text{MDA}(H_\xi)$, $\xi > 0$, $X \geq 0$, then:

$$k < \alpha = \frac{1}{\xi} \Rightarrow \mathbb{E}[X^k] < \infty$$

- (vi) *Remarks:*

- Distributions in the Fréchet MDA have tails that decay like power functions.
- Survival function (approximation): $\bar{H}_\xi(x) \approx (\xi x)^{-1/\xi}$
- The Fréchet MDA is the most important in practice.
- *Examples:* inverse gamma, Student-t, log-gamma, Cauchy, Pareto

- $\xi = 0$: **Gumbel MDA**

(\leadsto rather light-tailed, decays exponentially)

- (i) Suppose that F is twice differentiable on some interval (c, x_F) and further that $F' > 0$ and $F'' < 0$ on that interval. Then if:

$$\lim_{x \rightarrow x_F} \frac{(1 - F(x)) F''(x)}{(F'(x))^2} = -1$$

it holds that $F \in \text{MDA}(H_{\xi=0})$.

- (ii) *Gumbel CDF:*

$$\Gamma(x) = \exp(-e^{-x}), \quad x \in \mathbb{R}$$

- (iii) *Right endpoint of F :* both $x_{H_0} < \infty$ and $x_{H_0} = \infty$ possible

- (iv) *Moments:* All moments exist, and if $X \sim F$ is non-negative, then all moments are finite.

- (v) *Remarks:*

- $\text{MDA}(H_0)$ contains all densities whose tails decay roughly exponentially (light-tailed), but the tails can be quite different (up to moderately heavy).

– *Examples:* normal, log-normal, exponential, gamma, standard Weibull, generalized hyperbolic (except Student-t)

- $\xi < 0$: **Weibull MDA** (\leadsto short-tailed)

- (i) For $\xi < 0$:

$$F \in \text{MDA}(H_\xi) \iff \bar{F}\left(x_F - \frac{1}{x}\right) = x^{\frac{1}{\xi}} L(x) \text{ and } x_F < \infty$$

for some $L \in \mathcal{R}_0$ (i.e. L a slowly varying function at ∞).

- (ii) *Weibull CDF:*

$$\Psi_\alpha(x) = \begin{cases} 1 & : x > 0 \\ \exp(-(-x)^\alpha) & : x \leq 0 \end{cases} \quad \alpha > 0$$

with shape parameter $\xi = \frac{1}{\alpha}$.

- (iii) *Normalizing sequences:* $a_n = x_F - F^{\leftarrow}(1 - \frac{1}{n})$ and $b_n = x_F$.

- (iv) *Right endpoint of F :* $x_{H_\xi} < \infty$

- (v) *Moments:* All moments exist, and if $X \sim F$ is non-negative, then all moments are finite.

- (vi) *Examples:* beta, uniform (with $\alpha = \beta = 1$)

Fisher-Tippett-Gnedenko Theorem

- If $F \in \text{MDA}(H)$ for some non-degenerate H , then H must be of GEV type, i.e. $H = H_\xi$ for some $\xi \in \mathbb{R}$. Thus:

$$F \in \text{MDA}(H) \Rightarrow H \stackrel{\text{type}}{\sim} H_\xi$$

- *Remarks:*

- Two CDFs F and G are of the same type if $\exists a > 0$ and $b \in \mathbb{R}$ s.t. $F(x) = G\left(\frac{x-b}{a}\right)$.

- *Interpretation:* If location-scale transformed maxima converge in distribution to a non-degenerate limit, the limiting distribution must be GEV distribution.

- We can always choose normalizing sequences $(a_n) > 0$, (b_n) s.t. H_ξ appears in canonical form.

- All commonly encountered continuous distributions are in the MDA of a GEV distribution.

Examples of distributions per MDA Consider:

Fréchet MDA $\xi > 0$	Gumbel MDA $\xi = 0$	Weibull MDA $\xi < 0$
Student-t	normal \mathcal{N}	beta
Pareto	log-normal	(uniform \mathcal{U})
inverse gamma	exponential	
log-gamma	gamma	
Cauchy	standard Weibull	
F distribution	generalized hyperbol.	

5.2 Threshold exceedances/Peaks-over-threshold

Generalized Pareto distribution (GPD)

- The CDF of the **generalized Pareto distribution (GPD)** is:

$$G_{\xi, \beta}(x) = \begin{cases} 1 - \left(1 + \frac{\xi x}{\beta}\right)^{-\frac{1}{\xi}} & \text{if } \xi \neq 0 \\ 1 - \exp\left(-\frac{x}{\beta}\right) & \text{if } \xi = 0 \end{cases}$$

where $\beta > 0$ is the *scale parameter*, ξ is the *shape parameter*, and the support is:

- $x \geq 0$ when $\xi \geq 0$,
- and $x \in \left[0, -\frac{\beta}{\xi}\right]$ when $\xi < 0$.

- The PDF of the GDP is given by:

$$g_{\xi, \beta}(x) = \begin{cases} \frac{1}{\beta} \left(1 + \frac{\xi x}{\beta}\right)^{-\frac{1}{\xi}-1} & \text{if } \xi \neq 0 \\ \frac{1}{\beta} \exp\left(-\frac{x}{\beta}\right) & \text{if } \xi = 0 \end{cases}$$

and the support is defined as for the CDF.

- Special cases of the *shape parameter* ξ :

- $\xi > 0$: Pareto $\left(\frac{1}{\xi}, \frac{\beta}{\xi}\right)$
- $\xi = 0$: Exp $\left(\frac{1}{\beta}\right)$
- $\xi < 0$: short-tailed Pareto type II distribution

- **Remarks:**

- The larger ξ , the heavier tailed is $G_{\xi, \beta}$.
- *Maximum domain of attraction*: $G_{\xi, \beta} \in \text{MDA}(H_{\xi})$
- The GPD is the canonical CDF for modelling excess losses over high u .

Excess distribution over u , mean excess function

- Let $X \sim F$. The **excess distribution over the threshold u** is:

$$F_u(x) = \mathbb{P}[X - u \leq x | X > u] = \frac{F(x+u) - F(u)}{1 - F(u)}, \quad x \in [0, x_F - u]$$

F_u describes the distribution of the excess loss over u , given that u is exceeded.

- If $\mathbb{E}[|X|] < \infty$, the **mean excess function** is the mean w.r.t. F_u :

$$e(u) = \mathbb{E}[X - u | X > u] = \frac{1}{\mathbb{P}[X > u]} \int_u^{x_F} (x - u) f(x) dx = \frac{1}{1 - F(u)} \int_u^{x_F} (1 - F(x)) dx$$

- The **sample mean excess function** is given by:

$$e_n(u) = \frac{\sum_{i=1}^n (x_i - u) \mathbb{I}_{\{x_i > u\}}}{\sum_{i=1}^n \mathbb{I}_{\{x_i > u\}}}$$

- The **sample mean excess plot** consists of the points $\{x_{(i)}, e_n(x_{(i)}) : 2 \leq i \leq n\}$, where $x_{(i)}$ denotes the i^{th} order statistic.

- If a distribution $F \in \text{MDA}(H_{\xi})$, then its mean excess plot has a slope equal to $\frac{\xi}{1-\xi}$.
- The mean excess plot of the Pareto distribution is expected to be *fast linear*.

- For the **GPD**, i.e. if $F = G_{\xi, \beta}$, it holds that:

$$F_u(x) = G_{\xi, \beta(u)}(x), \quad \beta(u) = \beta + \xi u$$

$$e(u) = \frac{\beta(u)}{1 - \xi} = \frac{\beta + \xi u}{1 - \xi}$$

$$\text{where } \text{supp } u = \begin{cases} 0 \leq u < \infty & \text{if } 0 \leq \xi < 1 \\ 0 \leq u \leq -\frac{\beta}{\xi} & \text{if } \xi < 0 \end{cases}$$

Note that the mean excess function $e(u)$ is linear in the threshold u , which is a characterizing property of the GPD.

- For continuous $X \sim F$ with $\mathbb{E}[|X|] < \infty$, the following formula holds for *expected shortfall*:

$$\text{ES}_{\alpha}(X) = \text{VaR}_{\alpha}(X) + e(\text{VaR}_{\alpha}(X)), \quad \alpha \in (0, 1)$$

Pickands-Balkema-de Haan Theorem

- Let F be a general distribution function and denote by x_F the right endpoint of F .

- Then we can find a (positive-measurable) function $\beta(u) > 0$ s.t.

$$\lim_{x \rightarrow x_F} \sup_{0 \leq x < x_F - u} |F_u(x) - G_{\xi, \beta(u)}(x)| = 0$$

$$\iff F \in \text{MDA}(H_{\xi}), \xi \in \mathbb{R}$$

Peaks-over-threshold (POT) approach

- Since the block-maxima-methods (BMM) is wasteful of data (i.e. only the maxima of large blocks are used), it has been largely superseded in practice by methods based on *threshold exceedances*.
- The **peaks-over-threshold (POT) approach** (threshold exceedances) uses all data above a designated high threshold u .

The method is as follows:

- Given losses $X_1, \dots, X_n \sim F \in \text{MDA}(H_{\xi})$, $\xi \in \mathbb{R}$, let:
 - $N_u = |\{i \in \{1, \dots, n\} : X_i > u\}|$: number of exceedances over the (given) threshold u
 - $\tilde{X}_1, \dots, \tilde{X}_{N_u}$: exceedances
 - $Y_k = \tilde{X}_k - u$, $k \in \{1, \dots, N_u\}$, the corresponding excesses
- If Y_1, \dots, Y_{N_u} are i.i.d. and (roughly) distributed as $G_{\xi, \beta}$, then the *log-likelihood* is given by:

$$\ell(\xi, \beta; Y_1, \dots, Y_{N_u}) = \sum_{k=1}^{N_u} \log g_{\xi, \beta}(Y_k)$$

$$= -N_u \log(\beta) - \left(1 + \frac{1}{\xi}\right) \sum_{k=1}^{N_u} \log \left(1 + \frac{\xi Y_k}{\beta}\right)$$

Then, maximize w.r.t. $\beta > 0$ and $1 + \frac{\xi Y_k}{\beta} > 0, \forall k \in \{1, \dots, N_u\}$.

Methods for choosing the threshold u

- Plot the mean excess function $e(u) := \mathbb{E}[X - u | X > u]$ against u . Then look for the lowest value u_0 of u s.t. $e(u)$ is linear for $u > u_0$.
- Fix u and estimate the shape parameter $\xi = \xi(u)$. Do this for various values of u . Plot $\xi(u)$ against u and look for the lowest value u_0 of u s.t. $\xi(u)$ is approximately constant for $u > u_0$.

Smith estimator

- The **Smith estimator** is the *tail estimator* defined as:

$$\hat{F}(x) = \frac{N_u}{n} \left(1 + \xi \frac{x - u}{\hat{\beta}}\right)^{-\frac{1}{\xi}}, \quad x \geq u$$

$$= \underbrace{\frac{1}{n} \sum_{i=1}^n \mathbb{I}_{\{X_i > u\}}}_{\hat{F}_u(u)} \underbrace{\left(1 + \xi \frac{x - u}{\hat{\beta}}\right)^{-\frac{1}{\xi}}}_{\hat{F}(x-u) \approx 1 - G_{\xi, \beta}(x-u)}$$

- The Smith estimator faces a *bias-variance tradeoff*: If u is increased, the bias of parametrically estimating $\hat{F}_u(x - u)$ decreases, but the variance of it and the nonparametrically estimated $\hat{F}(u)$ increases.

Hill estimator

- Assume $F \in \text{MDA}(H_\xi)$, $\xi > 0$, so that $\bar{F}(x) = x^{-\alpha} L(x)$, $\alpha > 0$.
- The standard form of the **Hill estimator** of the *tail index* α is:

$$\hat{\alpha}_{k,n}^{(H)} = \left(\frac{1}{k} \sum_{i=1}^k \log X_{i,n} - \log X_{k,n} \right)^{-1}, \quad 2 \leq k \leq n$$

with k sufficiently small.

- Choosing k : Find a small k where the *Hill plot* stabilizes.
- **Semi-parametric Hill tail estimator**

$$\hat{\bar{F}}(x) = \frac{k}{n} \left(\frac{x}{X_{k,n}} \right)^{-\hat{\alpha}_{k,n}^{(H)}}, \quad x \geq X_{k,n}$$

- **Semi-parametric Hill VaR estimator**

$$\widehat{\text{VaR}}_\alpha(X) = \left(\frac{n}{k} (1 - \alpha) \right)^{-\frac{1}{\hat{\alpha}_{k,n}^{(H)}}} X_{k,n}, \quad \alpha \geq F(u) \approx 1 - \frac{k}{n}$$

- The semi-parametric **Hill ES estimator** is for $\alpha_{k,n}^{(H)} > 1$, $\alpha \geq F(u) \approx 1 - \frac{k}{n}$:

$$\begin{aligned} \widehat{\text{ES}}_\alpha(X) &= \frac{\left(\frac{n}{k} \right)^{-\frac{1}{\hat{\alpha}_{k,n}^{(H)}}} X_{k,n}}{1 - \alpha} \int_\alpha^1 (1 - z)^{-\frac{1}{\hat{\alpha}_{k,n}^{(H)}}} dz \\ &= \frac{\hat{\alpha}_{k,n}^{(H)}}{\hat{\alpha}_{k,n}^{(H)} - 1} \widehat{\text{VaR}}_\alpha(X) \end{aligned}$$

- **Observations from simulation study:**

- The empirical $\text{VaR}_{0.99}$ estimator has a negative bias.
- The Hill $\text{VaR}_{0.99}$ estimator has a negative bias for small k but a rapidly growing positive bias for larger k .
- The GDP $\text{VaR}_{0.99}$ estimator has a positive bias which grows much more slowly.
- The GDP $\text{VaR}_{0.99}$ estimator attains lowest MSE for a value of k around 100, but the MSE is a very robust choice of k (because of the slow growth of the bias) \rightarrow choice of u is less critical.
- The Hill $\text{VaR}_{0.99}$ estimator performs well for $20 \leq k \leq 75$ but then deteriorates rapidly.
- Both EVT methods outperform the empirical quantile estimator.

6 Multivariate models

6.1 Basics of multivariate modelling

Joint and marginal distributions

- Let $\mathbf{X} = (X_1, \dots, X_d) : \Omega \rightarrow \mathbb{R}^d$ be a d -dimensional *random vector*.
- The **(joint) distribution function (CDF)** F of \mathbf{X} is:

$$F_{\mathbf{X}}(\mathbf{x}) = \mathbb{P}[\mathbf{X} \leq \mathbf{x}] = \mathbb{P}[X_1 \leq x_1, \dots, X_d \leq x_d]$$

- The j^{th} **marginal** or **marginal CDF** F_j of \mathbf{X} is:

$$F_j(x_j) = \mathbb{P}[X_j \leq x_j] = F(\infty, \dots, \infty, x_j, \infty, \dots, \infty)$$

(interpreted as limit).

k -dimensional margins: For $\mathbf{X}_1 = (X_1, \dots, X_k)^\top$ and $\mathbf{X}_2 = (X_{k+1}, \dots, X_d)^\top$, the marginal CDF of \mathbf{X}_1 is:

$$F_{\mathbf{X}_1}(\mathbf{x}_1) = \mathbb{P}(\mathbf{X}_1 \leq \mathbf{x}_1) = F(x_1, \dots, x_k, \infty, \dots, \infty)$$

- F is **absolutely continuous** if:

$$F(\mathbf{x}) = \int_{(-\infty, \mathbf{x}]} f(\mathbf{z}) d\mathbf{z}$$

for some $f \geq 0$ known as the **(joint) density of \mathbf{X} (or F)**.

The j^{th} marginal CDF F_j is absolutely continuous if $F_j(x) = \int_{-\infty}^x f_j(z) dz$ for some $f_j \geq 0$ known as the density of X_j (or F_j).

- **Survival function \bar{F} of \mathbf{X} :**

$$\bar{F}_{\mathbf{X}}(\mathbf{x}) = \mathbb{P}[\mathbf{X} > \mathbf{x}] = \mathbb{P}[X_1 > x_1, \dots, X_d > x_d]$$

with corresponding j^{th} *marginal survival function* \bar{F}_j :

$$\bar{F}_j(x_j) = \mathbb{P}[X_j > x_j] = \bar{F}(\infty, \dots, \infty, x_j, \infty, \dots, \infty)$$

Note that in general: $\bar{F}(\mathbf{x}) \neq 1 - F(\bar{\mathbf{x}})$ (unless $d = 1$, i.e. in the univariate case).

- **Remarks:**

- Existence of a *joint density* \rightarrow existence of *marginal densities* for all k -dimensional marginals, $1 \leq k \leq d - 1$. The converse is false in general.
- **Discrete case:** replace integrals by sums IOT obtain similar formulas (the notion of densities is then replaced by *probability mass functions*)

Conditional distributions and independence

- Let $\mathbf{X} = (\mathbf{X}_1^\top, \mathbf{X}_2^\top)^\top \sim F$.
- The **conditional CDF of \mathbf{X}_2 given $\mathbf{X}_1 = \mathbf{x}_1$** is:

$$\begin{aligned} F_{\mathbf{X}_2|\mathbf{X}_1}(\mathbf{x}_2|\mathbf{x}_1) &= \mathbb{P}[\mathbf{X}_2 \leq \mathbf{x}_2 | \mathbf{X}_1 = \mathbf{x}_1] \\ &= \mathbb{E}[\mathbb{I}_{\{\mathbf{X}_2 \leq \mathbf{x}_2\}} | \mathbf{X}_1 = \mathbf{x}_1] \end{aligned}$$

- The **conditional PDF of \mathbf{X}_2 given $\mathbf{X}_1 = \mathbf{x}_1$** is:

$$f_{\mathbf{X}_2|\mathbf{X}_1}(\mathbf{x}_2|\mathbf{x}_1) = \frac{f(\mathbf{x}_1, \mathbf{x}_2)}{f_{\mathbf{X}_1}(\mathbf{x}_1)}$$

- Useful identities:

$$\begin{aligned} F(\mathbf{x}) &= \int_{(-\infty, \mathbf{x}_1]} F_{\mathbf{X}_2|\mathbf{X}_1}(\mathbf{x}_2|\mathbf{z}) dF_{\mathbf{X}_1}(\mathbf{z}) \\ F_{\mathbf{X}_2|\mathbf{X}_1}(\mathbf{x}_2|\mathbf{x}_1) &= \int_{-\infty}^{x_{k+1}} \dots \int_{-\infty}^{x_d} f_{\mathbf{X}_2|\mathbf{X}_1}(z_{k+1}, \dots, z_d | \mathbf{x}_1) dz_{k+1} \dots dz_d \end{aligned}$$

- **Characteristic function (CF)**

$$\varphi_{\mathbf{X}} = \mathbb{E} \left[e^{i\mathbf{t}^\top \mathbf{X}} \right], \quad \mathbf{t} \in \mathbb{R}^d$$

- **Independence:**

- \mathbf{X}, \mathbf{Y} are *independent*

\iff (CDFs)

$$F(\mathbf{X}, \mathbf{Y}) = F_{\mathbf{X}}(\mathbf{x}) F_{\mathbf{Y}}(\mathbf{y}), \quad \forall \mathbf{x}, \mathbf{y}$$

\iff (PDFs)

$$f(\mathbf{x}, \mathbf{y}) = f_{\mathbf{X}}(\mathbf{x}) f_{\mathbf{Y}}(\mathbf{y}), \quad \forall \mathbf{x}, \mathbf{y} \text{ (if PDFs exist!)} \\ \text{(in this case: } f_{\mathbf{Y}|\mathbf{X}}(\mathbf{y}|\mathbf{x}) = f_{\mathbf{Y}}(\mathbf{y}))$$

\iff (Characteristic functions)

$$\varphi_{(\mathbf{X}, \mathbf{Y})}(\mathbf{x}, \mathbf{y}) = \varphi_{\mathbf{X}}(\mathbf{x}) \cdot \varphi_{\mathbf{Y}}(\mathbf{y}) \\ \text{(in this case: } \varphi_{\mathbf{X}+\mathbf{Y}}(\mathbf{z}) = \varphi_{\mathbf{X}}(\mathbf{z}) \cdot \varphi_{\mathbf{Y}}(\mathbf{z}))$$

If two RVs \mathbf{X}, \mathbf{Y} are independent, then:

$$\mathbb{E}[\mathbf{XY}] = \mathbb{E}[\mathbf{X}]\mathbb{E}[\mathbf{Y}], \quad \text{Cov}[\mathbf{XY}] = \mathbb{E}[\mathbf{XY}] - \mathbb{E}[\mathbf{X}]\mathbb{E}[\mathbf{Y}]$$

Remark: The converse does not hold in general. But independence can be rejected if it can be shown that one of these properties does not hold.

- Similarly, the components X_1, \dots, X_d of \mathbf{X} are *mutually independent*

$$\iff F(\mathbf{x}) = \prod_{j=1}^d F_j(x_j) \quad \iff f(\mathbf{x}) = \prod_{j=1}^d f_j(x_j)$$

$$\iff \varphi_{\mathbf{X}}(\mathbf{t}) = \prod_{j=1}^d \phi_{X_j}(t_j)$$

$\forall \mathbf{x}$ or $\forall \mathbf{t}$, respectively, and if the PDF f exists.

- Two RVs $\mathbf{X}, \mathbf{Y} \in \mathbb{R}^d$ are **equal in distribution** \iff

$$\mathbf{a}^\top \mathbf{X} \stackrel{d}{=} \mathbf{a}^\top \mathbf{Y}, \quad \forall \mathbf{a} \in \mathbb{R}^d$$

Moments and characteristic function

- If $\mathbb{E}[|X_j|] < \infty, \forall j$, the **mean vector** of \mathbf{X} is:

$$\mathbb{E}[\mathbf{X}] = (\mathbb{E}[X_1], \dots, \mathbb{E}[X_d])^\top$$

It holds that:

$$X_1, \dots, X_d \text{ independent} \Rightarrow \mathbb{E}[X_1 \cdots X_d] = \prod_{j=1}^d \mathbb{E}[X_j]$$

- If $\mathbb{E}[X_j^2] < \infty, \forall j$, the **covariance matrix** of \mathbf{X} is:

$$\Sigma := \text{Cov}[\mathbf{X}] = \mathbb{E}[(\mathbf{X} - \mathbb{E}[\mathbf{X}])(\mathbf{X} - \mathbb{E}[\mathbf{X}])^\top]$$

The $(i, j)^{\text{th}}$ element of Σ is:

$$\begin{aligned} \sigma_{ij} &= \Sigma_{ij} = \text{Cov}[X_i, X_j] \\ &= \mathbb{E}[(X_i - \mathbb{E}[X_i])(X_j - \mathbb{E}[X_j])] = \mathbb{E}[X_i X_j] - \mathbb{E}[X_i]\mathbb{E}[X_j] \\ \sigma_{ii} &= \text{Var}[X_i] \end{aligned}$$

Note that: $\mathbf{X}_1, \mathbf{X}_2$ independent $\Rightarrow \text{Cov}[X_1, X_2] = 0$.

$\text{Cov}[X_1, X_2] = 0 \Rightarrow$ independence holds true only for the bivariate normal distribution.

A counter-example is the bivariate Student-t distribution with $\nu > 2$ degrees of freedom.

- If $\mathbb{E}[X_j^2] < \infty, \forall j$, the **correlation matrix** of \mathbf{X} is given by $\text{Corr}[\mathbf{X}]$ with the $(i, j)^{\text{th}}$ element:

$$\text{Corr}[X_i, X_j] = \frac{\text{Cov}[X_i, X_j]}{\sqrt{\text{Var}[X_i] \text{Var}[X_j]}}$$

which is in $[-1, 1]$.

Note that: $\text{Corr}[X_i, X_j] = \pm 1$ iff $X_j = aX_i + b$.

- **Properties:** ($\forall \mathbf{A} \in \mathbb{R}^{k \times d}, \mathbf{b} \in \mathbb{R}^k, \mathbf{a} \in \mathbb{R}^d$)

$$- \mathbb{E}[\mathbf{A}\mathbf{X} + \mathbf{b}] = \mathbf{A}\mathbb{E}[\mathbf{X}] + \mathbf{b} = \mathbf{A}\boldsymbol{\mu} + \mathbf{b}$$

$$- \text{Cov}[\mathbf{A}\mathbf{X} + \mathbf{b}] = \mathbf{A} \text{Cov}[\mathbf{X}] \mathbf{A}^\top$$

If $k = 1$ ($\mathbf{A} = \mathbf{a}^\top$), then:

$$\mathbf{a}^\top \Sigma \mathbf{a} = \text{Cov}[\mathbf{a}^\top \mathbf{X}] = \text{Var}[\mathbf{a}^\top \mathbf{X}] \geq 0$$

- It holds that:

A symmetric matrix Σ is a covariance matrix

$\iff \Sigma$ is *positive semidefinite*.

If Σ is a positive definite matrix, then Σ is *invertible*.

- The **Colesky decomposition** is:

$$\Sigma = \mathbf{A}\mathbf{A}^\top$$

for a lower triangular matrix (*Cholesky factor*) \mathbf{A} with $A_{ii} > 0, \forall j$.

Standard estimators of covariance and correlation

- Assume:

$\mathbf{X}_1, \dots, \mathbf{X}_n \sim F$, serially uncorrelated and with:

$$\boldsymbol{\mu} := \mathbb{E}[\mathbf{X}_1], \quad \Sigma := \text{Cov}[\mathbf{X}_1], \quad P := \text{Corr}[\mathbf{X}_1]$$

- **Non-parametric method-of-moments-like estimators:**

- *sample mean:*

$$\bar{\mathbf{X}} = \frac{1}{n} \sum_{i=1}^n \mathbf{X}_i$$

with $\text{Cov}[\bar{\mathbf{X}}] = \frac{1}{n} \Sigma$ and which is clearly unbiased.

- *sample covariance matrix:*

$$S = \frac{1}{n} \sum_{i=1}^n (\mathbf{X}_i - \bar{\mathbf{X}})(\mathbf{X}_i - \bar{\mathbf{X}})^\top \quad (\text{biased})$$

$$\begin{aligned} S_n &= \frac{1}{n-1} \sum_{i=1}^n (\mathbf{X}_i - \bar{\mathbf{X}})(\mathbf{X}_i - \bar{\mathbf{X}})^\top \quad (\text{unbiased}) \\ &= \frac{n}{n-1} S \end{aligned}$$

S_n is unbiased since it can be shown that $\mathbb{E}[S_n] = \Sigma$.

- *sample correlation matrix:*

$$R = (R_{ij}), \quad R_{ij} = \frac{S_{ij}}{\sqrt{S_{ii} S_{jj}}}$$

Multivariate normal distribution

- $\mathbf{X} = (X_1, \dots, X_d)$ has a **multivariate normal (or Gaussian) distribution** if:

$$\mathbf{X} = \boldsymbol{\mu} + \mathbf{A}\mathbf{Z}$$

where $\mathbf{Z} = (Z_1, \dots, Z_k)$, $Z_l \sim \mathcal{N}(0, 1)$, $\mathbf{A} \in \mathbb{R}^{d \times k}$, $\boldsymbol{\mu} \in \mathbb{R}^d$.

Its *mean* and *covariance* are:

$$\mathbb{E}[\mathbf{X}] = \boldsymbol{\mu}, \quad \text{Cov}[\mathbf{X}] = \mathbf{A}\mathbf{A}^\top =: \Sigma$$

- **Characteristic function:**

$$\varphi_{\mathbf{X}}(\mathbf{t}) = \mathbb{E} \left[e^{i\mathbf{t}^\top \mathbf{X}} \right] = \exp \left(i\mathbf{t}^\top \boldsymbol{\mu} - \frac{1}{2} \mathbf{t}^\top \Sigma \mathbf{t} \right)$$

- For $\mathbf{X} \sim \mathcal{N}_d(\boldsymbol{\mu}, \Sigma)$ with $\text{rank } \mathbf{A} = d = k$ ($\Rightarrow \Sigma$ positive definite, invertible), the **density** of \mathbf{X} is:

$$f_{\mathbf{X}}(\mathbf{x}) = \frac{1}{(2\pi)^{d/2} \sqrt{\det \Sigma}} \exp \left(-\frac{1}{2} (\mathbf{x} - \boldsymbol{\mu})^\top \Sigma^{-1} (\mathbf{x} - \boldsymbol{\mu}) \right)$$

- **Independence of components:** The following statements are equivalent:

\iff The components of $\mathbf{X} \sim \mathcal{N}_d(\boldsymbol{\mu}, \Sigma)$ are *mutually independent*.

$\iff \Sigma$ is *diagonal*.

\iff The components of \mathbf{X} are *uncorrelated* (and joint normally distributed).

- **Transformations** for $\mathbf{X} \sim \mathcal{N}_d(\boldsymbol{\mu}, \Sigma)$:

- In general: $\mathbf{a}^\top \mathbf{X} \sim \mathcal{N}(\mathbf{a}^\top \boldsymbol{\mu}, \mathbf{a}^\top \Sigma \mathbf{a})$, $\forall \mathbf{a} \in \mathbb{R}^d$

- *Margins:* $X_j \sim \mathcal{N}(\mu_j, \sigma_j^2)$

- *Sums:* $\sum_{j=1}^d X_j \sim \mathcal{N}(\sum_{j=1}^d \mu_j, \sum_{i,j} \sigma_{ij})$

- *Linear combinations:* for $\mathbf{B} \in \mathbb{R}^{k \times d}$, $\mathbf{b} \in \mathbb{R}^k$, it holds:

$$\mathbf{B}\mathbf{X} + \mathbf{b} \sim \mathcal{N}_k(\mathbf{B}\boldsymbol{\mu} + \mathbf{b}, \mathbf{B}\Sigma\mathbf{B}^\top)$$

- *Convolutions:* for an independent $\mathbf{Y} \sim \mathcal{N}_d(\tilde{\boldsymbol{\mu}}, \tilde{\Sigma})$, it holds:

$$\mathbf{X} + \mathbf{Y} \sim \mathcal{N}_d(\boldsymbol{\mu} + \tilde{\boldsymbol{\mu}}, \Sigma + \tilde{\Sigma})$$

- Convolution for dependent $\mathbf{X}_1, \mathbf{X}_2$:

Assume the following known joint distribution:

$$\mathbf{X} = \begin{pmatrix} \mathbf{X}_1 \\ \mathbf{X}_2 \end{pmatrix} \sim \mathcal{N}_{2d} \left(\boldsymbol{\mu} = \begin{pmatrix} \boldsymbol{\mu}_1 \\ \boldsymbol{\mu}_2 \end{pmatrix}, \Sigma = \begin{pmatrix} \Sigma_1 & \Sigma_{12} \\ \Sigma_{12} & \Sigma_2 \end{pmatrix} \right)$$

Note that \mathbf{X} has margins $\mathbf{X}_i \sim \mathcal{N}_d(\boldsymbol{\mu}_i, \Sigma_i), i = 1, 2$ and that Σ_{12} describes the dependence structure between $\mathbf{X}_1, \mathbf{X}_2$.

Then the sum $\mathbf{X}_1 + \mathbf{X}_2$ can be expressed as:

$$\mathbf{X}_1 + \mathbf{X}_2 = \mathbf{A}^\top \mathbf{X}, \quad \mathbf{A} = \begin{pmatrix} \mathbb{I}_d \\ \mathbb{I}_d \end{pmatrix}, \mathbb{I}_d \in \mathbb{R}^{d \times d}$$

which has the following distribution:

$$\begin{aligned} \mathbf{X}_1 + \mathbf{X}_2 &\sim \mathcal{N}_d \left(\mathbf{A}^\top \boldsymbol{\mu}, \mathbf{A}^\top \Sigma \mathbf{A} \right) \\ &\sim \mathcal{N}_d(\boldsymbol{\mu}_1 + \boldsymbol{\mu}_2, \Sigma_1 + 2\Sigma_{12} + \Sigma_2) \\ &\stackrel{\text{in general}}{\sim} \mathcal{N}_d(\boldsymbol{\mu}_1 + \boldsymbol{\mu}_2, \Sigma_1 + \Sigma_2) \end{aligned}$$

- **Sampling of $\mathcal{N}_d(\boldsymbol{\mu}, \Sigma)$**

(i) Compute the *Cholesky factor* \mathbf{A} of Σ .

(ii) Generate $\mathbf{Z} = (Z_1, \dots, Z_d)$ with independent $Z_j \sim \mathcal{N}(0, 1)$.

(iii) Return $\mathbf{X} = \boldsymbol{\mu} + \mathbf{A}\mathbf{Z}$.

6.2 Normal mixture distributions

Multivariate normal variance mixtures

- The random vector \mathbf{X} has a **(multivariate) normal variance mixture distribution** if:

$$\mathbf{X} \stackrel{d}{=} \boldsymbol{\mu} + \sqrt{W} \mathbf{A} \mathbf{Z}$$

where $\mathbf{Z} \sim \mathcal{N}_k(\mathbf{0}, \mathbb{I}_k)$, $W \geq 0$ is a RV independent of \mathbf{Z} , $\mathbf{A} \in \mathbb{R}^{d \times k}$ with the *scale matrix* $\Sigma = \mathbf{A} \mathbf{A}^\top$, and $\boldsymbol{\mu} \in \mathbb{R}^d$ the *location vector*.

Notation: $\mathbf{X} \sim M_k(\boldsymbol{\mu}, \Sigma, \hat{F}_W)$

■ Remarks:

- Note that:
 $(\mathbf{X}|W = w) = \boldsymbol{\mu} + \sqrt{w} \mathbf{A} \mathbf{Z} \sim N_d(\boldsymbol{\mu}, w\Sigma)$
- W can be interpreted as a shock affecting the variances of all risk factors.

■ Properties:

Let $\mathbf{X} = \boldsymbol{\mu} + \sqrt{W} \mathbf{A} \mathbf{Z}$ and $\mathbf{Y} = \boldsymbol{\mu} + \mathbf{A} \mathbf{Z}$. Assume that $\text{rank}(\mathbf{A}) = d \leq k$ and that Σ is positive definite.

- If $\mathbb{E}[\sqrt{W}] < \infty$, then $\mathbb{E}[\mathbf{X}] = \boldsymbol{\mu} = \mathbb{E}[\mathbf{Y}]$.
- If $\mathbb{E}[W] < \infty$, then:
 $\text{Cov}[\mathbf{X}] = \mathbb{E}[W]\Sigma = \mathbb{E}[W]\mathbf{A}\mathbf{A}^\top \neq \Sigma = \text{Cov}[\mathbf{Y}]$ in general
- **Linear combinations:**
For $\mathbf{X} \sim M_d(\boldsymbol{\mu}, \Sigma, \hat{F}_W)$ and $\mathbf{Y} = \mathbf{B}\mathbf{X} + \mathbf{b}$, where $\mathbf{B} \in \mathbb{R}^{k \times d}$ and $\mathbf{b} \in \mathbb{R}^k$, we have:
 $\mathbf{Y} \sim M_k(\mathbf{B}\boldsymbol{\mu} + \mathbf{b}, \mathbf{B}\Sigma\mathbf{B}^\top, \hat{F}_W)$
If $\mathbf{a} \in \mathbb{R}^d$, then $\mathbf{a}^\top \mathbf{X} \sim M_1(\mathbf{a}^\top \boldsymbol{\mu}, \mathbf{a}^\top \Sigma \mathbf{a}, \hat{F}_W)$.
- If $\mathbb{E}[W] < \infty$, then $\text{Corr}[\mathbf{X}] = \text{Corr}[\mathbf{Y}]$.

■ Independence:

Let $\mathbf{X} = \boldsymbol{\mu} + \sqrt{W} \mathbf{Z}$ with $\mathbb{E}[W] < \infty$ (uncorrelated normal variance mixture). Then:

X_i and X_j are independent $\iff W$ is a.s. constant (i.e. $\mathbf{X} \sim \mathcal{N}_d(\boldsymbol{\mu}, W\mathbb{I}_k)$)

■ Characteristic function (CF):

The CF of a multivariate normal variance mixture is:

$$\varphi_{\mathbf{X}}(\mathbf{t}) = \exp(i\mathbf{t}^\top \boldsymbol{\mu}) \mathbb{E} \left[\exp \left(-W \frac{1}{2} \mathbf{t}^\top \Sigma \mathbf{t} \right) \right]$$

■ Density:

If Σ is positive definite and $\mathbb{P}[W = 0] = 0$, the density of \mathbf{X} is:

$$f_{\mathbf{X}}(\mathbf{x}) = \int_0^\infty \frac{1}{(2\pi)^{d/2} w^{d/2} |\Sigma|^{1/2}} \cdot \exp \left(-\frac{(\mathbf{x} - \boldsymbol{\mu})^\top \Sigma^{-1} (\mathbf{x} - \boldsymbol{\mu})}{2w} \right) dF_W(w)$$

- The **Laplace-Stieltjes (LS) transform** of F_W is:

$$\hat{F}_W(\theta) := \mathbb{E}[\exp(-\theta W)] = \int_0^\infty e^{-\theta w} dF_W(w)$$

■ Sampling:

- Generate $\mathbf{Z} \sim \mathcal{N}_d(\mathbf{0}, \mathbb{I}_d)$.
- Generate $W \sim F_W$, independent of \mathbf{Z} .
- Compute the Cholesky factor \mathbf{A} (s.t. $\mathbf{A} \mathbf{A}^\top = \Sigma$).
- Return $\mathbf{X} = \boldsymbol{\mu} + \sqrt{W} \mathbf{A} \mathbf{Z}$.

■ Examples:

- Student-t distribution:
The stochastic representation of a RV $\mathbf{X} = t_d(\nu, \boldsymbol{\mu}, \Sigma)$ is:

$$\mathbf{X} \stackrel{d}{=} \boldsymbol{\mu} + \Sigma^{1/2} \mathbf{Y}$$

where $\mathbf{Y} \sim t_d(\nu, \mathbf{0}, \mathbb{I}_d)$ and $\Sigma^{1/2}$ the Cholesky factor of Σ . This can also be written in terms of a normal variance mixture:

$$\mathbf{X} \stackrel{d}{=} \boldsymbol{\mu} + \sqrt{W} \Sigma^{1/2} \mathbf{Z} \stackrel{d}{=} \boldsymbol{\mu} + \sqrt{\frac{\nu}{V}} \Sigma^{1/2} \mathbf{Z}$$

where $\mathbf{Z} \sim \mathcal{N}_d(\mathbf{0}, \mathbb{I}_d)$, $W \sim \text{Ig}(\frac{\nu}{2}, \frac{\nu}{2})$ (inverse gamma distribution) and $V \sim \chi_\nu^2$ independent of \mathbf{Z} .

Then: $\text{Cov}[\mathbf{X}] = \mathbb{E}[W]\Sigma$.

- **Remark:** Since normal variance mixtures are elliptical distributions, VaR is subadditive and thus a coherent risk measure for normal variance mixtures.

Normal mean-variance mixtures

- \mathbf{X} has **(multivariate) normal mean-variance mixture distribution** if

$$\mathbf{X} = \mathbf{m}(W) + \sqrt{W} \mathbf{A} \mathbf{Z}$$

where $\mathbf{Z} \sim \mathcal{N}_k(\mathbf{0}, \mathbb{I}_k)$, $W \geq 0$ a scalar random variable (independent of \mathbf{Z}), $\mathbf{A} \in \mathbb{R}^{d \times k}$ a matrix of constants, $\mathbf{m} : [0, \infty) \rightarrow \mathbb{R}^d$ a measurable function.

- **Remark:** Normal mean-variance mixtures add *skewness*. In general, these distributions are no longer elliptical.

■ Examples:

- Let $\mathbf{m}(W) = \boldsymbol{\mu} + W\boldsymbol{\gamma}$. It then holds that $\mathbb{E}[\mathbf{X}|W] = \boldsymbol{\mu} + W\boldsymbol{\gamma}$ and $\text{Cov}[\mathbf{X}|W] = W\Sigma$.
We then have:

$$\mathbb{E}[\mathbf{X}] = \boldsymbol{\mu} + \mathbb{E}[W]\boldsymbol{\gamma} \quad \text{if } \mathbb{E}[W] < \infty$$

$$\text{Cov}[\mathbf{X}] = \mathbb{E}[W]\Sigma + \text{Var}[W]\boldsymbol{\gamma}\boldsymbol{\gamma}^\top \quad \text{if } \mathbb{E}[W^2] < \infty$$

- Let $\mathbf{X} \stackrel{d}{=} \boldsymbol{\mu} + \sqrt{W} \mathbf{Z}$ in $d = 2$.
Then $\text{Cov}[X_1, X_2] = \mathbb{E}[W Z_1 Z_2]$.
In this case, if Z_1, Z_2 are uncorrelated and W is a constant, then X_1, X_2 are independent. If W is not a constant, then X_1, X_2 are dependent through W .

6.3 Spherical and elliptical distributions

Spherical distribution

- A random vector $\mathbf{Y} = (Y_1, \dots, Y_d)$ has a **spherical distribution** if for every *orthogonal* $\mathbf{U} \in \mathbb{R}^{d \times d}$ (i.e. with $\mathbf{U}\mathbf{U}^\top = \mathbf{U}^\top \mathbf{U} = \mathbb{I}_d$):

$$\mathbf{Y} \stackrel{d}{=} \mathbf{U} \mathbf{Y} \quad (\text{distr. invariant under rotations and reflections})$$

■ Characterisation of spherical distributions:

Let $\|\mathbf{t}\| = \|\mathbf{t}\|_{L^2}$, $\mathbf{t} \in \mathbb{R}^d$. The following are equivalent:

$$\iff \mathbf{Y} \text{ is spherical (notation: } \mathbf{Y} \sim S_d(\psi))$$

$$\iff \exists \text{ a characteristic generator } \psi : [0, \infty) \rightarrow \mathbb{R}, \text{ s.t.}$$

$$\varphi_{\mathbf{Y}}(\mathbf{t}) = \mathbb{E}[e^{i\mathbf{t}^\top \mathbf{Y}}] = \psi(\|\mathbf{t}\|^2), \forall \mathbf{t} \in \mathbb{R}^d.$$

$$\iff \text{For every } \mathbf{a} \in \mathbb{R}^d, \mathbf{a}^\top \mathbf{Y} = \|\mathbf{a}\| Y_1 \text{ (linear combination of the same type).}$$

Remark: The third statement implies that there is *subadditivity* of VaR_α for jointly elliptical losses.

■ Stochastic representation:

$\mathbf{Y} \sim S_d(\psi)$ iff $\mathbf{Y} \stackrel{d}{=} R\mathbf{S}$ for an independent *radial part* $R \geq 0$ and $\mathbf{S} \sim \mathcal{U}(\{\mathbf{x} \in \mathbb{R}^d : \|\mathbf{x}\| = 1\})$.

Elliptical distribution

- A random vector $\mathbf{X} = (X_1, \dots, X_d)$ has an **elliptical distribution** if:

$$\mathbf{X} \stackrel{d}{=} \boldsymbol{\mu} + \mathbf{A} \mathbf{Y}, \quad (\text{multivariate affine transformation})$$

where $\mathbf{Y} \sim S_k(\psi)$, $\mathbf{A} \in \mathbb{R}^{d \times k}$ (*scale matrix* $\Sigma = \mathbf{A} \mathbf{A}^\top$), and *location vector* $\boldsymbol{\mu} \in \mathbb{R}^d$.

If Σ is positive definite with Cholesky factor \mathbf{A} , then $\mathbf{X} \sim E_d(\boldsymbol{\mu}, \Sigma, \psi)$ iff $\mathbf{Y} = \mathbf{A}^{-1}(\mathbf{X} - \boldsymbol{\mu}) \sim S_d(\psi)$.

Remark: normal variance mixture distributions are (all) elliptical.

■ Characteristic function:

$$\varphi_{\mathbf{X}}(\mathbf{t}) = \mathbb{E}[e^{i\mathbf{t}^\top \mathbf{X}}] = e^{i\mathbf{t}^\top \boldsymbol{\mu}} \psi(\mathbf{t}^\top \Sigma \mathbf{t}).$$

Notation: $\mathbf{X} \sim E_d(\boldsymbol{\mu}, \Sigma, \psi)$.

■ Stochastic representation:

$\mathbf{X} \stackrel{d}{=} \boldsymbol{\mu} + \mathbf{R} \mathbf{A} \mathbf{S}$, with \mathbf{R} and \mathbf{S} as above.

■ Properties:

Density:

Let Σ be positive definite and $\mathbf{Y} \sim S_d(\psi)$ have density generator g .

Then $\mathbf{X} = \mu + A\mathbf{Y}$ has density:

$$f_{\mathbf{X}}(\mathbf{x}) = \frac{1}{\sqrt{\det \Sigma}} g((\mathbf{x} - \mu)^\top \Sigma^{-1} (\mathbf{x} - \mu))$$

This density is constant on ellipsoids, i.e. on the sets:

$$\{\mathbf{x} \in \mathbb{R}^d : (\mathbf{x} - \mu)^\top \Sigma (\mathbf{x} - \mu) = \text{const.}\}$$

Linear combinations:

For $\mathbf{X} \sim E_d(\mu, \Sigma, \psi)$, $B \in \mathbb{R}^{k \times d}$ and $\mathbf{b} \in \mathbb{R}^k$:

$$B\mathbf{X} + \mathbf{b} \sim E_k(B\mu + \mathbf{b}, B\Sigma B^\top, \psi)$$

For $\mathbf{a} \in \mathbb{R}^d$:

$$\mathbf{a}^\top \mathbf{X} \sim E_1(\mathbf{a}^\top \mu, \mathbf{a}^\top \Sigma \mathbf{a}, \psi)$$

Thus, all *marginal distributions* are of the same type.

Marginal densities:

Margins of elliptical distributions are elliptical since for $\mathbf{X} = (\mathbf{X}_1^\top, \mathbf{X}_2^\top)^\top \sim E_d(\mu, \Sigma, \psi)$ satisfies $\mathbf{X}_1 \sim E_k(\mu_1, \Sigma_{11}, \psi)$ and $\mathbf{X}_2 \sim E_{d-k}(\mu_2, \Sigma_{22}, \psi)$.

Conditional distributions:

Conditional distributions of elliptical distributions are elliptical. *Conditional correlations* remain invariant.

Quadratic forms:

It holds that $(\mathbf{X} - \mu)^\top \Sigma^{-1} (\mathbf{X} - \mu) = R^2$.

If $\mathbf{X} \sim \mathcal{N}_d(\mu, \Sigma)$, then $R^2 \sim \chi_d^2$.

If $\mathbf{X} \sim t_d(\nu, \mu, \Sigma)$, then $\frac{R^2}{d} \sim F(d, \nu)$.

Convolutions:

Let $\mathbf{X} \sim E_d(\mu, \Sigma, \psi)$ and $\mathbf{Y} \sim E_d(\tilde{\mu}, c\Sigma, \tilde{\psi})$ be independent. Then $a\mathbf{X} + b\mathbf{Y}$ is elliptically distributed for $a, b \in \mathbb{R}, c > 0$.

- **Remark:** VaR is subadditive and thus a coherent risk measure for elliptical distributions.

7 Copulas and dependence

7.1 Copulas

Copulas

Reasoning/Motivation:

F " = " marginal dfs F_1, \dots, F_d " + " dependence structure C

Advantages:

- Most natural in a *static distributional context* (i.e. no time dependence, e.g. on residuals of an ARMA-GARCH model).
- Copulas allow us to understand and study *dependence independently of the margins*.
- Copulas allow for a *bottom-up approach* to multivariate model building (e.g. to construct tailored F).

Copulas:

A copula C is a CDF with $\mathcal{U}(0, 1)$ margins.

$C : [0, 1]^d \rightarrow [0, 1]$ is a copula iff:

- C is *grounded*, i.e.
 $C(u_1, \dots, u_d) = 0$ if $u_j = 0$ for at least one $j \in \{1, \dots, d\}$.
- C has *standard uniform univariate margins*, i.e.
 $C(1, \dots, 1, u_j, 1, \dots, 1) = u_j$ for all $u_j \in [0, 1]$ and $j \in \{1, \dots, d\}$.
- C is *d-increasing*, i.e.
 C assigns non-negative mass to all non-empty hypercubes in $[0, 1]^d$.
Equivalently (if existent): density $c(\mathbf{u}) \geq 0$ for all $\mathbf{u} \in (0, 1)^d$.

Transformations

Probability transformation

Let $X \sim F$, F continuous. Then $F(X) \sim \mathcal{U}(0, 1)$.

Quantile transformation

Let $U \sim \mathcal{U}(0, 1)$ and F be any CDF. Then $X = F^{\leftarrow}(U) \sim F$.

- **Remark:** Probability and quantile transformations are the key to all applications involving copulas. They allow us to go from \mathbb{R}^d to $[0, 1]^d$ and back.

Sklar's Theorem

CDFs of copulas

- For any CDF F with margins F_1, \dots, F_d , there exists a copula C s.t.

$$F(x_1, \dots, x_d) = C(F_1(x_1), \dots, F_d(x_d)), \quad \mathbf{x} \in \mathbb{R}^d$$

C is uniquely defined on $\prod_{j=1}^d \text{ran } F_j$.

- Define the RV $\mathbf{X} = (X_1, \dots, X_d)$ s.t. $\mathbf{X} \sim F$ (with margins F_i). Then C is given by:

$$C(u_1, \dots, u_d) = F(F_1^{\leftarrow}(u_1), \dots, F_d^{\leftarrow}(u_d)), \quad \mathbf{u} \in \prod_{j=1}^d \text{ran } F_j$$

$$= \mathbb{P}[F_1(X_1) \leq u_1, \dots, F_d(X_d) \leq u_d]$$

- Conversely, given any copula C and univariate CDFs F_1, \dots, F_d , F as defined above is a CDF with margins F_1, \dots, F_d .
- *Interpretation:*
The *first part* allows one to decompose any CDF F into its margins and a copula.
This (together with the invariance principle) allows one to study dependence independently of the margins via the margin-free $\mathbf{U} = (F_1(X_1), \dots, F_d(X_d))$ instead of $\mathbf{X} = (X_1, \dots, X_d)$ (since they both have the same copula).
 \leadsto statistical applications, e.g. parameter estimation or goodness-of-fit.
The *second part* allows one to construct flexible multivariate distributions for particular applications.

PDFs of copulas

- If the CDF F_j has PDF f_j , $j \in \{1, \dots, d\}$, and the CDF C has PDF c , then the PDF f of F satisfies:

$$f(\mathbf{x}) = c(F_1(x_1), \dots, F_d(x_d)) \prod_{j=1}^d f_j(x_j)$$

$$\log f(\mathbf{x}) = \log c(F_1(x_1), \dots, F_d(x_d)) + \sum_{j=1}^d \log f_j(x_j)$$

and we can recover the copula's PDF c via:

$$c(\mathbf{u}) = \frac{f(F_1^{-1}(u_1), \dots, F_d^{-1}(u_d))}{f_1(F_1^{-1}(u_1)) \cdots f_d(F_d^{-1}(u_d))}$$

- Note that *not* all copulas have a PDF.

Invariance principle

- Let $X_j \sim F_j$, F_j continuous, $j \in \{1, \dots, d\}$. Then:

$$\mathbf{X} \sim F \text{ has copula } C \iff (F_1(X_1), \dots, F_d(X_d)) \sim C$$

- Let $\mathbf{X} \sim F$ with continuous margins F_1, \dots, F_d and copula C . If T_j is strictly increasing ($T_j \uparrow$) on $\text{ran } X_j$ for all j , then $(T_1(X_1), \dots, T_d(X_d))$ has again copula C .

6.4 Dimension reduction techniques

Remark: not part of the exam.

Fréchet-Hoeffding bounds

- Define the copulas W and M as follows:

$$W(\mathbf{u}) = \left(\sum_{j=1}^d u_j - d + 1 \right)^+ \quad (\text{countermonotone})$$

$$M(\mathbf{u}) = \min_{1 \leq j \leq d} \{u_j\} \quad (\text{comonotone})$$

Then the following holds:

(i) (Fréchet-Hoeffding bounds)

For any d -dimensional copula C :

$$W(\mathbf{u}) \leq C(\mathbf{u}) \leq M(\mathbf{u}), \quad \mathbf{u} \in [0, 1]^d$$

(ii) W is a copula iff $d = 2$.

(iii) M is a copula $\forall d \geq 2$.

- The following bounds for any CDF F can be derived from the Fréchet-Hoeffding bounds:

$$\left(\sum_{j=1}^d F_j(x_j) - d + 1 \right)^+ \leq F(\mathbf{x}) \leq \min_{1 \leq j \leq d} \{F_j(x_j)\}$$

■ Remarks:

- It holds for the uniform distribution $U \sim \mathcal{U}(0, 1)$:

$$(U, \dots, U) \sim M, \quad (U, 1 - U) \sim W$$

- The Fréchet-Hoeffding bounds correspond to *perfect dependence*, i.e. negative for W , positive for M .

Examples of copulas

■ Fundamental copulas

The **independence copula** is $\Pi(\mathbf{u}) = \prod_{j=1}^d u_j$ since:

$$C(F_1(x_1), \dots, F_d(x_d)) = F(\mathbf{x}) = \prod_{j=1}^d F_j(x_j)$$

$$\iff C(\mathbf{u}) = \Pi(\mathbf{u})$$

Therefore, X_1, \dots, X_d are independent iff their copula is Π .

- The Fréchet-Hoeffding bound W is the *countermonotonicity copula*, which is the CDF of $(U, 1 - U)$.
If X_1, X_2 are perfectly *negatively* dependent (i.e. X_2 is a strictly decreasing function in X_1), then their copula is W .
- The Fréchet-Hoeffding bound M is the *comonotonicity copula*, which is the CDF of (U, \dots, U) .
If X_1, \dots, X_d are perfectly *positively* dependent (i.e. X_2, \dots, X_d are a.s. a strictly increasing functions in X_1), then their copula is M .

■ Implicit copulas (elliptical copulas)

The **elliptical copulas** are implicit copulas arising from elliptical distributions via Sklar's Theorem.

– Gauss copulas

- (i) Consider (w.l.o.g.) $\mathbf{X} \sim \mathcal{N}_d(\mathbf{0}, P)$. The **Gauss copula** (family) is:

$$C_P^{\text{Ga}}(\mathbf{u}) = \mathbb{P}[\Phi(X_1) \leq u_1, \dots, \Phi(X_d) \leq u_d]$$

$$= \Phi_P(\Phi^{-1}(u_1), \dots, \Phi^{-1}(u_d))$$

where Φ_P is the CDF of $\mathcal{N}_d(\mathbf{0}, P)$ and Φ the CDF of $\mathcal{N}(0, 1)$.

- (ii) The *PDF* of $C(\mathbf{u})$ is:

$$c(\mathbf{u}) = \frac{f(F_1^{-1}(u_1), \dots, F_d^{-1}(u_d))}{\prod_{j=1}^d f_j(F_j^{-1}(u_j))}$$

In particular, the PDF of C_P^{Ga} is:

$$c_P^{\text{Ga}} = \frac{1}{\sqrt{\det P}} \exp\left(-\frac{1}{2} \mathbf{x}^\top (P^{-1} - \mathbb{I}_d) \mathbf{x}\right)$$

where $\mathbf{x} = (\Phi^{-1}(u_1), \dots, \Phi^{-1}(u_d))$.

- (iii) *Special cases:*

$$P = \mathbb{I}_d \Rightarrow C = \Pi, \quad P = \mathbb{J}_d = \mathbf{1}\mathbf{1}^\top \Rightarrow C = M$$

$$d = 2 \text{ and } \rho = P_{12} = -1 \Rightarrow C = W$$

– t copulas

- (i) Consider (w.l.o.g.) $\mathbf{X} \sim t_d(\nu, \mathbf{0}, P)$. The **t copula** (family) is:

$$C_{\nu, P}^t(\mathbf{u}) = \mathbb{P}[t_\nu(X_1) \leq u_1, \dots, t_\nu(X_d) \leq u_d]$$

$$= t_{\nu, P}(t_\nu^{-1}(u_1), \dots, t_\nu^{-1}(u_d))$$

where $t_{\nu, P}$ is the CDF of $t_d(\nu, \mathbf{0}, P)$ and t_ν the CDF of the univariate t distribution with ν degrees of freedom.

- (ii) The *PDF* of $C_{\nu, P}^t(\mathbf{u})$ is:

$$c_{\nu, P}^t(\mathbf{u}) = \frac{\Gamma((\nu + d)/2)}{\Gamma(\nu/2)\sqrt{\det P}} \left(\frac{\Gamma(\nu/2)}{\Gamma((\nu + 1)/2)} \right)^d$$

$$\cdot \frac{(1 + \mathbf{x}^\top P^{-1} \mathbf{x}/\nu)^{-(\nu + d)/2}}{\prod_{j=1}^d (1 + x_j^2/\nu)^{-(\nu + 1)/2}}$$

- (iii) *Special cases:*

$$P = \mathbb{J}_d = \mathbf{1}\mathbf{1}^\top \Rightarrow C = M$$

$$d = 2 \text{ and } \rho = P_{12} = -1 \Rightarrow C = W$$

However: $P = \mathbb{I}_d \Rightarrow C \neq \Pi$ (unless $\nu = \infty \Rightarrow C_{\nu, P}^t = C_P^{\text{Ga}}$)

- *Remark:* Elliptical copulas are symmetric/exchangeable.

– Sampling of implicit copulas:

- (i) Sample $\mathbf{X} \sim F$, where F is a density function with continuous margins F_1, \dots, F_d
- (ii) Return $\mathbf{U} = (F_1(X_1), \dots, F_d(X_d))$ (probability transformation)

■ Explicit copulas (Archimedean copulas)

- The **Archimedean copulas** are copulas of the form:

$$C(\mathbf{u}) = \psi(\psi^{-1}(u_1) + \dots + \psi^{-1}(u_d)), \quad \mathbf{u} \in [0, 1]^d$$

with the **Archimedean generator** ψ where:

- (i) $\psi : [0, \infty) \rightarrow [0, 1]$;
- (ii) $\psi(0) = 1$ and $\psi(\infty) = \lim_{t \rightarrow \infty} \psi(t) = 0$;
- (iii) ψ is continuous and strictly decreasing on: $(0, \inf\{x : \psi(x) = 0\})$;
- (iv) $\psi^{-1}(0) = \inf\{t : \psi(t) = 0\}$ by convention.

The set of all generators ψ is denoted by Ψ . If $\psi(t) > 0$, $t \in [0, \infty)$, then we call ψ *strict*.
 ψ can be interpreted as the *Laplace transform* of a non-negative RV $V \sim G$.

- *Examples:*

(i) Clayton copula

$$\psi(t) = (1 + t)^{-1/\theta}, \quad t \in [0, \infty), \quad \theta \in (0, \infty)$$

$$\Rightarrow C_\theta^C(\mathbf{u}) = (u_1^{-\theta} + \dots + u_d^{-\theta} - d + 1)^{-1/\theta}$$

Limits: For $\theta \downarrow 0$, $C \rightarrow \Pi$, and for $\theta \uparrow \infty$, $C \rightarrow M$.

(ii) Gumbel copula

$$\psi(t) = \exp(-t^{1/\theta}), \quad t \in [0, \infty), \quad \theta \in [1, \infty)$$

$$\Rightarrow C_\theta^G(\mathbf{u}) = \exp\left(-\left((-\log u_1)^\theta + \dots + (-\log u_d)^\theta\right)^{1/\theta}\right)$$

Limits: For $\theta = 1$, $C \rightarrow \Pi$, and for $\theta \rightarrow \infty$, $C \rightarrow M$.

- *Remark:* Archimedean copulas are symmetric/exchangeable.

– Simulation of Archimedean copulas (Marshall and Olkin):

If ψ is the Laplace transform of a non-negative RV $V \sim G$:

- (i) Generate $V \sim G$ (CDF corresponding to ψ);
- (ii) generate $E_1, \dots, E_d \stackrel{\text{i.i.d.}}{\sim} \text{Exp}(1)$, independent of V ;
- (iii) return:

$$\mathbf{U} = \left(\psi\left(\frac{E_1}{V}\right), \dots, \psi\left(\frac{E_d}{V}\right) \right)^\top$$

(conditional independence)

– **Simulation of Archimedean copulas (using \mathcal{U}):**

If ψ is the Laplace transform of a non-negative RV $V \sim G$:

- (i) Generate $V \sim G$;
- (ii) generate $X_1, \dots, X_d \stackrel{\text{i.i.d.}}{\sim} \mathcal{U}(0, 1)$;
- (iii) return:

$$\mathbf{U} = \left(\hat{G} \left(-\log \frac{X_1}{V} \right), \dots, \hat{G} \left(-\log \frac{X_d}{V} \right) \right)^\top$$

Survival copulas

- If $\mathbf{U} \sim C$, then the **survival copula** of C is given by $\mathbf{1} - \mathbf{U} \sim \hat{C}$. The survival copula \hat{C} can be expressed as:

$$\hat{C}(\mathbf{u}) = \sum_{J \subseteq \{1, \dots, d\}} (-1)^{|J|} C \left((1 - u_1)^{\mathbb{1}_{J(1)}}, \dots, (1 - u_d)^{\mathbb{1}_{J(d)}} \right)$$

For $d = 2$:

$$\hat{C}(u_1, u_2) = u_1 + u_2 - 1 + C(1 - u_1, 1 - u_2)$$

If C admits a *PDF*, then $\hat{c}(\mathbf{u}) = c(\mathbf{1} - \mathbf{u})$.

It holds for the *tail dependence coefficients*:

$$\lambda_u^{\hat{C}} = \lambda_l^C, \quad \lambda_l^{\hat{C}} = \lambda_u^C$$

- **Sklar's theorem** for survival copulas:

$$\bar{F}(\mathbf{x}) = \hat{C}(\bar{F}_1(x_1), \dots, \bar{F}_d(x_d)), \quad \mathbf{x} \in \mathbb{R}^d$$

where $F(\mathbf{x}) = \mathbb{P}[\mathbf{X} > \mathbf{x}]$ with corresponding marginal survival functions $\bar{F}_1, \dots, \bar{F}_d$ (with $\bar{F}_j(x) = \mathbb{P}[X_j > x]$).

- **Radially symmetric copulas**

- If $\hat{C} = C$, then C is called **radially symmetric**.
radially symmetric copulas: e.g. W, Π, M , Gauss copulas and t -copulas
radially symmetric copula family: e.g. elliptical copulas
- **Tail dependence coefficients**: $\lambda_u = \lambda_l =: \lambda$
- If X_j is symmetrically distributed about a_j , $j \in \{1, \dots, d\}$, then \mathbf{X} is radially symmetric about \mathbf{a} iff $C = \hat{C}$.

- **Remark**: Survival copulas combine marginal survival functions to joint survival functions.

Note that while \hat{C} is a CDF, \bar{F} and $\bar{F}_1, \dots, \bar{F}_d$ are *not* CDFs!

Exchangeability

- \mathbf{X} is **exchangeable** if:

$$(X_1, \dots, X_d) = (X_{\pi(1)}, \dots, X_{\pi(d)})$$

for any permutation $(\pi(1), \dots, \pi(d))$ of $(1, \dots, d)$.

- A *copula* C is **exchangeable** if it is the CDF of an exchangeable \mathbf{U} with $\mathcal{U}(0, 1)$ margins.
This holds iff $C(u_1, \dots, u_d) = C(u_{\pi(1)}, \dots, u_{\pi(d)})$, i.e. if c is *symmetric*.

- **Remarks**:

- Exchangeable/symmetric copulas are useful for approximate modelling of homogeneous portfolios.
- Examples: Archimedean copulas, elliptical copulas (e.g. Gauss, t) for equicorrelated P (i.e. $P = \rho \mathbb{J}_d + (1 - \rho) \mathbb{I}_d$ for $\rho \geq \frac{-1}{d-1}$, in particular $d = 2$).

7.2 Dependence concepts and measures

Measures of association/dependence are scalar measures which summarize the dependence in terms of a single number.

Perfect dependence

- **Counter/comonotonicity**

- X_1, X_2 are **countermonotone** if (X_1, X_2) has copula W .
- X_1, \dots, X_d are **comonotone** if (X_1, \dots, X_d) has copula M .
equivalently: X_1, \dots, X_d are comonotone if $(L_1, \dots, L_d) \stackrel{d}{=} (f_1(Z), \dots, f_d(Z))$ for some RV Z and non-decreasing transformations f_1, \dots, f_d .

- **Perfect dependence**

- $X_2 = T(X_1)$ a.s. with decreasing $T(x) = F_2^{\leftarrow}(1 - F_1(x))$ (countermonotone)
 $\iff C(u_1, u_2) = W(u_1, u_2), u_1, u_2 \in [0, 1]$.
- $X_j = T_j(X_1)$ a.s. with increasing $T_j(x) = F_j^{\leftarrow}(F_1(x))$, $j \in \{2, \dots, d\}$ (comonotone)
 $\iff C(\mathbf{u}) = M(\mathbf{u}), \mathbf{u} \in [0, 1]^d$.

- **Comonotone additivity**

Let $\alpha \in (0, 1)$ and $X_j \sim F_j$, $j \in \{1, \dots, d\}$, be *comonotone*.
Then $F_{X_1 + \dots + X_d}^{\leftarrow}(\alpha) = F_1^{\leftarrow}(\alpha) + \dots + F_d^{\leftarrow}(\alpha)$.

Linear correlation

- The linear correlation coefficient $\rho(X_1, X_2)$ does not always exist, i.e. if the second moment does not exist ($\mathbb{E}[X_i^2] = \infty$ or not defined).

ρ is *not* an exclusive copula property, i.e. ρ depends on both the copula and the marginals.

- For two RVs X_1, X_2 with $\mathbb{E}[X_j^2] < \infty$, $j \in \{1, 2\}$, the **linear (or Pearson's) correlation coefficient** $\rho = \text{Corr}[X_1, X_2]$ is:

$$\begin{aligned} \rho(X_1, X_2) &= \frac{\text{Cov}[X_1, X_2]}{\sqrt{\text{Var}[X_1] \text{Var}[X_2]}} \\ &= \frac{\mathbb{E}[X_1 X_2] - \mathbb{E}[X_1] \mathbb{E}[X_2]}{\sqrt{\mathbb{E}[(X_1 - \mathbb{E}[X_1])^2] \mathbb{E}[(X_2 - \mathbb{E}[X_2])^2]}} \\ &= \frac{\mathbb{E}[(X_1 - \mathbb{E}[X_1])(X_2 - \mathbb{E}[X_2])]}{\sqrt{\mathbb{E}[(X_1 - \mathbb{E}[X_1])^2] \mathbb{E}[(X_2 - \mathbb{E}[X_2])^2]}} \end{aligned}$$

Remarks: For two RVs X, Y :

- $\text{Var}[X, Y]$ is maximal if the linear correlation $\text{Corr}[X, Y]$ is maximal.
(since $\text{Var}[X + Y] = \text{Var}[X] + \text{Var}[Y] + 2 \text{Corr}[X, Y] \sqrt{\text{Var}[X] \text{Var}[Y]}$ and variance is always non-negative)
- $\text{VaR}_\alpha(X, Y)$ is *not* maximal if the linear correlation $\text{Corr}[X, Y]$ is maximal (see previous example).
- $\rho(X, Y)$ is maximal if X, Y are *comonotone*, i.e. if X, Y have copula $M(u, v) = \min(u, v)$.

- The **Hoeffding's identity** is:

$$\text{Cov}[X, Y] = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} F(x, y) - F_X(x) F_Y(y) dx dy$$

- **Properties**:

Let X_1, X_2 be two RVs with $\mathbb{E}[X_j^2] < \infty$, $j \in \{1, 2\}$.

- It always holds that $|\rho| \leq 1$.
 $|\rho| = 1 \iff \exists a \in \mathbb{R} \setminus \{0\}, b \in \mathbb{R} \text{ with } X_2 = aX_1 + b$.
It holds that $a > 0 \Rightarrow \rho = +1$ and $a < 0 \Rightarrow \rho = -1$.

- X_1 and X_2 are *independent* $\nRightarrow \rho = 0$.

The converse is in general *not* true.

An example where zero linear correlation also implies independence is the *multivariate normal distribution*.

- ρ is *invariant* under strictly increasing *linear* transformations on $\text{ran } X_1 \times \text{ran } X_2$ but *not invariant* under strictly increasing functions in general.

- **Correlation fallacies**

- **Fallacy 1**: F_1, F_2 and ρ uniquely determine F .
This is true for *bivariate elliptical distributions*, but wrong in general.
- **Fallacy 2**: Given F_1, F_2 , any $\rho \in [-1, 1]$ is attainable.
This is true for *elliptically distributed* (X_1, X_2) with $\mathbb{E}[R^2] < \infty$ (since then $\text{Corr}[\mathbf{X}] = P$), but wrong in general.
i.e. if F_1 and F_2 are not of the same type (no linearity), $\rho(X_1, X_2) = 1$ is not attainable.

- **Fallacy 3:** ρ maximal (i.e. $C = M$) $\Rightarrow \text{VaR}_\alpha(X_1 + X_2)$ maximal.

This is true if (X_1, X_2) are *elliptically distributed*.

- **Remark:** Increasing linear transformations of margins keep linear correlation unaffected.

Rank correlation coefficients

- Both rank correlation coefficients ρ_τ, ρ_S are *copula properties*, and are thus invariant under strictly increasing transformations of the underlying RVs.
 ρ_τ, ρ_S always exist for two continuous RVs X, Y .

■ Kendall's tau

Let $X \sim F_X, Y \sim F_Y$ with F_X, F_Y continuous. Let (X', Y') be an independent copy of (X, Y) . Then **Kendall's tau** is:

$$\rho_\tau(X, Y) = \mathbb{E}[\text{sign}((X - X')(Y - Y'))] \\ = \mathbb{P}[(X - X')(Y - Y') > 0] - \mathbb{P}[(X - X')(Y - Y') < 0]$$

If (X, Y) has copula $C(u, v)$, then:

$$\rho_\tau(X, Y) = 4 \int_0^1 \int_0^1 C(u, v) dC(u, v) - 1$$

Remarks:

- Kendall's tau is the the probability of *concordance* minus the probability of *discordance*.
- For any given marginal distributions, Kendall's tau can reach any value in $[-1, 1]$, depending on the chosen copula.
The *bounds* are given by:
 - (i) comonotone copula M : $\rho_\tau(M) = 1$
 - (ii) countermonotone copula W : $\rho_\tau(W) = -1$

■ Spearman's rho

Let $X \sim F_X, Y \sim F_Y$ with F_X, F_Y continuous.
Then **Spearman's rho** is:

$$\rho_S(X, Y) = \rho(F_X(X), F_Y(Y))$$

If F_X, F_Y have the copula C , then:

$$\rho_S(X, Y) = 12 \int_0^1 \int_0^1 (C(u, v) - uv) dudv \\ = 12 \int_0^1 \int_0^1 C(u, v) dudv - 3$$

An alternative definition uses the Pearson correlation coefficient applied to the ranks $\text{rank } X_i, \text{rank } Y_i$ of the samples $X_i, Y_i, i = 1, \dots, n$:

$$\rho_S(X, Y) = \rho_{\text{rank } X, \text{rank } Y} = \frac{\text{Cov}[\text{rank } X, \text{rank } Y]}{\sigma_{\text{rank } X} \sigma_{\text{rank } Y}}$$

Remarks: (κ either $\kappa = \rho_\tau$ or $\kappa = \rho_S$)

- In general, $\kappa = 0$ does *not* imply independence.
- The correlation fallacies 1 and 3 are *not* solved by replacing ρ by rank correlation coefficients κ .
- But correlation fallacy 2 (i.e. for F_X, F_Y , any $\rho \in [-1, 1]$ is attainable) is solved.

Coefficients of tail dependence

- **Goal:** Measure extremal dependence, i.e. dependence in the *joint tails*.

- The coefficients of tail dependence λ_l, λ_u are *copula properties*, and are thus invariant under strictly increasing transformations of the underlying RVs.

λ_l, λ_u are *not* defined for all pairs of RVs X_1, X_2 (limit!).

■ Tail dependence

Let $X_j \sim F_j, j \in \{1, 2\}$, be continuously distributed RVs. Provided that the limits exist, the following *equivalent* definitions of the **lower tail-dependence coefficient** λ_l and the **upper tail-dependence coefficient** λ_u of X_1 and X_2 exist:

- via the *inverse CDFs* F_1, F_2 :

$$\lambda_l = \lim_{q \downarrow 0} \mathbb{P}[X_2 \leq F_2^{\leftarrow}(q) \mid X_1 \leq F_1^{\leftarrow}(q)] \\ \lambda_u = \lim_{q \uparrow 1} \mathbb{P}[X_2 > F_2^{\leftarrow}(q) \mid X_1 > F_1^{\leftarrow}(q)]$$

(order of conditioning can be reversed)

- via VaR_α :

$$\lambda_l = \lim_{q \downarrow 0} \mathbb{P}[X_2 \leq \text{VaR}_q(X_2) \mid X_1 \leq \text{VaR}_q(X_1)] \\ \lambda_u = \lim_{q \uparrow 1} \mathbb{P}[X_2 > \text{VaR}_q(X_2) \mid X_1 > \text{VaR}_q(X_1)]$$

(order of conditioning can be reversed)

- via the *copula* C :

$$\lambda_l = \lim_{q \downarrow 0} \frac{C(q, q)}{q} = \lambda_l^{\hat{C}} \\ \lambda_u = 2 - \lim_{q \uparrow 1} \frac{1 - C(q, q)}{1 - q} = \lim_{q \uparrow 1} \frac{1 - 2q + C(q, q)}{1 - q} \\ = \lim_{q \downarrow 0} \frac{\hat{C}(q, q)}{q} = \lambda_l^{\hat{C}}$$

■ Asymptotic dependence/independence

- If $\lambda_l > 0$ / $\lambda_u > 0$, we say that there is **asymptotic dependence** in the lower/upper tail.
- If $\lambda_l = 0$ / $\lambda_u = 0$, we say that there is **asymptotic independence** in the lower/upper tail.

- **Remarks:**

- Tail dependence of the *survival copula* \hat{C} :
The upper/lower tail dependence coefficients of the survival copula \hat{C} are equal to the lower/upper tail dependence coefficient of the copula C , i.e.

$$\lambda_l^{\hat{C}} = \lambda_u^C, \quad \lambda_u^{\hat{C}} = \lambda_l^C$$

- λ_l, λ_u for the counter-/comonotone copulas W, M :

$$\lambda_l^W = \lambda_u^W = -1, \quad \lambda_l^M = \lambda_u^M = +1$$

- For all *radially symmetric* copulas (e.g. the bivariate $C_{\nu, P}^{\text{Ga}}$ and $C_{\nu, P}^t$), we have $\lambda_l = \lambda_u =: \lambda$.
- For *Archimedean copulas* with strict ψ , it holds e.g.
Clayton: $\lambda_l = 2^{-1/\theta}, \lambda_u = 0$
Gumbel: $\lambda_l = 0, \lambda_u = 2 - 2^{1-\theta}$
- *Comparison of Gauss copula with Student-t copula w.r.t. extreme values*:
For distributions with a Gauss copula, extreme values are independent (i.e. $\lambda_u = \lambda_l = 0$), while for distributions with a Student-t copula, extreme values are dependent.

7.3 Normal mixture copulas

Tail dependence (for normal mixture copulas)

■ Normal mixture copulas

The *normal mixture copulas* are the copulas of the *multivariate normal (mean-)variance mixtures*:

$$\mathbf{X} \stackrel{d}{=} \boldsymbol{\mu} + \sqrt{W} \mathbf{A} \mathbf{Z} \quad \text{or} \quad \mathbf{X} \stackrel{d}{=} \mathbf{m}(W) + \sqrt{W} \mathbf{A} \mathbf{Z}$$

■ Coefficients of tail dependence

Let (X_1, X_2) be distributed according to a normal variance mixture and assume (w.l.o.g.) that $\boldsymbol{\mu} = (0, 0)^\top$ and $\mathbf{A} \mathbf{A}^\top = \mathbf{P} = \begin{pmatrix} 1 & \rho \\ \rho & 1 \end{pmatrix}$.
In this case, $F_1 = F_2$ and C is *symmetric* and *radially symmetric*.
We thus obtain that:

$$\lambda = \lambda_l = 2 \lim_{x \downarrow -\infty} \mathbb{P}[X_2 \leq x \mid X_1 = x]$$

- **Remarks:**

- What drives tail dependence of normal variance mixtures is W .
If W has a power tail, we get tail dependence, otherwise not.
- *Covariance matrix* of normal (mean-)variance mixtures \mathbf{X} :

$$\text{Cov}[\mathbf{X}] = \mathbb{E}[W] \Sigma = \mathbb{E}[W] \mathbf{A} \mathbf{A}^\top$$

- **Examples:**

– Gauss copula

Considering the bivariate $\mathcal{N}(\mathbf{0}, P)$ density, one can show that $(X_2|X_1 = x) \sim \mathcal{N}(\rho x, 1 - \rho^2)$. This implies that $\lambda = \mathbb{I}_{\rho=1}$ (essentially no tail dependence).

– t copula

For $C_{\nu, P}$, one can show that:

$$(X_2|X_1 = x) \sim t_{\nu+1} \left(\rho x, \frac{(1 - \rho^2)(\nu + x^2)}{\nu + 1} \right)$$

$$\text{thus: } \mathbb{P}[X_2 \leq x|X_1 = x] = t_{\nu+1} \left(\frac{x - \rho x}{\sqrt{\frac{(1 - \rho^2)(\nu + x^2)}{\nu + 1}}} \right)$$

and hence:

$$\lambda = 2t_{\nu+1} \left(-\sqrt{\frac{(\nu + 1)(1 - \rho)}{1 + \rho}} \right) \quad (\text{tail dependence})$$

Rank correlations (for normal mixture copulas)

■ Spearman's rho for normal variance mixtures

Let $\mathbf{X} \sim M_2(\mathbf{0}, P, \hat{F}_W)$ with $\mathbb{P}[\mathbf{X} = \mathbf{0}] = 0$, $\rho = P_{12}$. Then:

$$\rho_S = \frac{6}{\pi} \mathbb{E} \left[\arcsin \left(\frac{W_\rho}{\sqrt{(W + \tilde{W})(W + \bar{W})}} \right) \right]$$

for $W, \tilde{W}, \bar{W} \sim F_W$ with Laplace-Stieltjes transform \hat{F}_W .
For *Gauss copulas*, $\rho_S = \frac{6}{\pi} \arcsin \left(\frac{\rho}{2} \right)$.

■ Kendall's tau for elliptical distributions

Let $\mathbf{X} \sim E_2(\mathbf{0}, P, \psi)$ with $\mathbb{P}[\mathbf{X} = \mathbf{0}] = 0$, $\rho = P_{12}$.

Then, $\rho_\tau = \frac{2}{\pi} \arcsin \rho$.

Skewed normal mixture copulas

■ Skewed normal mixture copulas are the copulas of normal mixture distributions which are *not elliptical*.

E.g. the skewed t copula $C_{\nu, P, \gamma}^t$ is the copula of a generalized hyperbolic distribution.

■ Remarks:

- It can be sampled as other implicit copulas.
- The main advantage of such a copula over $C_{\nu, P}^t$ is its *radial asymmetry* (e.g. for modelling $\lambda_l \neq \lambda_u$).

Grouped normal mixture copulas

■ Grouped normal mixture copulas are copulas which *attach together a set of normal mixture copulas*.

■ Example: a grouped t copula is the copula of:

$$\mathbf{X} = (\sqrt{W_1}Y_1, \dots, \sqrt{W_1}Y_{s_1}, \dots, \sqrt{W_S}Y_{s_1+\dots+s_{S-1}+1}, \dots, \sqrt{W_S}Y_d)$$

for $(W_1, \dots, W_S \sim M(\text{IG}(\frac{\nu_1}{2}, \frac{\nu_1}{2}), \dots, \text{IG}(\frac{\nu_S}{2}, \frac{\nu_S}{2})))$ and $\mathbf{Y} \sim \mathcal{N}_d(\mathbf{0}, P)$ (so $\mathbf{Y} = A\mathbf{Z}$).

Remarks:

- The marginals are t distributed, hence:

$$\mathbf{U} = (t_{\nu_1}(X_1), \dots, t_{\nu_1}(X_{s_1}), \dots, t_{\nu_S}(X_{s_1+\dots+s_{S-1}+1}), \dots, t_{\nu_S}(X_d))$$

follows a *grouped t copula*.

- It can be fitted with pairwise inversion of *Kendall's tau*.
- If $S = d$, grouped t copulas are also known as *generalized t copulas*.

7.4 Archimedean copulas

Bivariate Archimedean copulas

■ For $\psi \in \Psi$:

$C(u_1, u_2) = \psi(\psi^{-1}(u_1) + \psi^{-1}(u_2))$ is a copula
 $\iff \psi$ is convex.

■ For a strict and twice-continuously differentiable ψ , it holds that:

$$\rho_\tau = 1 - 4 \int_0^\infty t(\psi'(t))^2 dt = 1 + 4 \int_0^1 \frac{\psi^{-1}(t)}{(\psi^{-1}(t))'} dt$$

■ If ψ is strict, it holds that:

$$\lambda_l = 2 \lim_{t \rightarrow \infty} \frac{\psi'(2t)}{\psi'(t)}, \quad \lambda_u = 2 - 2 \lim_{t \downarrow 0} \frac{\psi'(2t)}{\psi'(t)}$$

Multivariate Archimedean copulas

■ ψ is *completely monotone (c.m.)* if $(-1)^k \psi^{(k)}(t) \geq 0$, $\forall t \in (0, \infty)$ and $\forall k \in \mathbb{N}_0$.

The set of all c.m. generators is denoted by Ψ_∞ .
Archimedean copulas with $\psi \in \Psi_\infty$ are called *LT-Archimedean copulas*.

■ Kimberling (1974)

If $\psi \in \Psi$, then $C(\mathbf{u}) = \psi \left(\sum_{j=1}^d \psi^{-1}(u_j) \right)$ is a copula

$\iff \psi \in \Psi_\infty$.

■ Bernstein (1928)

$\psi(0) = 1$, ψ c.m. $\iff \psi(t) = \mathbb{E}[\exp(-tV)]$ for $V \sim G$ with $V \geq 0$ and $G(0) = 0$.

■ Stochastic representation:

Let $\psi \in \Psi_\infty$ with $V \sim G$ s.t. $\hat{G} = \psi$ and let $E_1, \dots, E_d \sim \text{Exp}(1)$ be independent of V . Then:

- The *survival copula* of $\mathbf{X} = (\frac{E_1}{V}, \dots, \frac{E_d}{V})$ is *Archimedean* (with ψ).
- $\mathbf{U} = (\psi(X_1), \dots, \psi(X_d)) \sim C$ and the U_j 's are *conditionally independent* given V with $\mathbb{P}[U_j \leq u|V = v] = \exp(-v\psi^{-1}(u))$.

7.5 Fitting copulas to data

Setting

■ Let $\mathbf{X}, \mathbf{X}_1, \dots, \mathbf{X}_n$ be independent random vectors with CDF F , continuous margins F_1, \dots, F_d and copula C .

■ We assume that we have data $\mathbf{x}_1, \dots, \mathbf{x}_n$ interpreted as realizations of $\mathbf{X}_1, \dots, \mathbf{X}_n$.

■ Assume:

- $F_j = F_j(\cdot, \theta_{0,j})$ for some $\theta_{0,j} \in \Theta_j$, $j \in \{1, \dots, d\}$,
 $F_j(\cdot, \theta_{0,j})$ continuous $\forall \theta_j \in \Theta_j$.
- $C = C(\cdot, \theta_{0,C})$ for some $\theta_{0,C} \in \Theta_C$.

■ Thus, F has the true but unknown parameter vector $\theta_0 = (\theta'_{0,C}, \theta'_{0,1}, \dots, \theta'_{0,d})$ to be estimated.

Method-of-moments using rank correlation

■ Focus: $\theta_{0,C} = \theta_{0,C}$.

■ For $d = 2$, one can estimate $\theta_{0,C}$ by solving $\rho_\tau(\theta_C) = r_n^\tau$ w.r.t. θ_C , i.e.

$$\hat{\theta}_{n,C}^{\text{IKTE}} = \rho_\tau^{-1}(r_n^\tau) \quad (\text{inversion of Kendall's tau estimator (IKTE)})$$

where $\rho_\tau(\cdot)$ denotes Kendall's tau as a function in θ and r_n^τ is the sample version of Kendall's tau (computed from $\mathbf{X}_1, \dots, \mathbf{X}_n$ or pseudo-observations $\mathbf{U}_1, \dots, \mathbf{U}_n$).

■ The *standardized dispersion matrix* P for *elliptical copulas* can be estimated via *pairwise inversion of Kendall's tau*.

■ For *Gauss copulas*, it is preferable to use *Spearman's rho* based on:

$$\rho_S = \frac{6}{\pi} \arcsin \frac{\rho}{2} \approx \rho$$

The latter approximation error is comparably small, so that the matrix of pairwise sample versions of Spearman's rho is an estimator for the correlation matrix P .

■ For t copulas, \hat{P}_n^{IKTE} can be used to estimate P and then ν can be estimated via its MLE based on \hat{P}_n^{IKTE} .

Forming a pseudo-sample from the copula

- X_1, \dots, X_n (almost) never has $\mathcal{U}(0,1)$ margins. For applying a "copula approach" we thus need *pseudo-observations* from C .
- In general, we take $\hat{U}_i = (\hat{U}_{i1}, \dots, \hat{U}_{id}) = (\hat{F}_1(X_{i1}), \dots, \hat{F}_d(X_{id}))$, $i \in \{1, \dots, n\}$, where \hat{F}_j denotes an estimator of F_j .

- Possible choices of \hat{F}_j :

(i) *Non-parametric* estimators with *scaled empirical CDFs*, so:

$$\hat{U}_{ij} = \frac{n}{n+1} \hat{F}_{n,j}(X_{ij}) = \frac{R_{ij}}{n+1}$$

where R_{ij} denotes the *rank* of X_{ij} among all X_{1j}, \dots, X_{nj} .

- (ii) *Parametric* estimators (e.g. Student-t, Pareto), typically if n is small.
- (iii) *EVT-based* estimators; bodies are modelled empirically, tails semiparametrically via GDP.

Maximum likelihood estimation

- By Sklar's theorem, the PDF of F is given by:

$$f(\mathbf{x}, \boldsymbol{\theta}_0) = c(F_1(x_1; \boldsymbol{\theta}_{0,1}), \dots, F_d(x_d; \boldsymbol{\theta}_{0,d}); \boldsymbol{\theta}_{0,C}) \prod_{j=1}^d f_j(x_j; \boldsymbol{\theta}_{0,j})$$

from which the *log-likelihood* follows directly.

- The **maximum likelihood estimator (MLE)** of $\boldsymbol{\theta}_0$ is thus:

$$\boldsymbol{\theta}_n^{\text{MLE}} = \arg \sup_{\boldsymbol{\theta} \in \Theta} \ell(\boldsymbol{\theta}, \mathbf{X}_1, \dots, \mathbf{X}_n)$$

Probability distributions

Exponential distribution $X \sim \text{Exp}(\lambda)$, $\lambda > 0$

- PDF/CDF/Quantile:

$$f(x) = \lambda e^{-\lambda x}, \quad F(x) = 1 - e^{-\lambda x}, \quad F^{-1}(x) = -\frac{\log(1-x)}{\lambda}$$

- Mean/Standard deviation / Characteristic function:

$$\mu = \frac{1}{\lambda}, \quad \sigma = \frac{1}{\lambda}, \quad \varphi(t) = \frac{\lambda}{\lambda - it}$$

- Transformation:

$X \sim \text{Exp}(\lambda)$ has the same distribution as $\frac{E}{\lambda}$, $E \sim \text{Exp}(1)$.

Pareto distribution $X \sim \text{Pareto}(\lambda)$, $\lambda > 0$

- PDF/CDF/Quantile:

$$f(x) = \frac{\lambda}{x^{\lambda+1}}, \quad F(x) = 1 - \frac{1}{x^\lambda}, \quad F^{-1}(u) = \frac{1}{(1-u)^{\frac{1}{\lambda}}}$$

- Mean/Variance:

$$\mu = \begin{cases} \infty & \text{for } \lambda \leq 1 \\ \frac{\lambda}{\lambda-1} & \text{for } \lambda > 1 \end{cases}, \quad \sigma^2 = \begin{cases} \infty & \text{for } \lambda \in (0, 2] \\ \frac{\lambda}{(\lambda-1)^2(\lambda-2)} & \text{for } \lambda > 2 \end{cases}$$

- Characteristic function:

$$\varphi(t) = \lambda(-it)^\lambda \Gamma(-\lambda, -it)$$

- Relation to the exponential distribution:

If $X \sim \text{Pareto}(\alpha)$, then $Y = \log X \sim \text{Exp}(\alpha)$.

Equivalently, if $Y \sim \text{Exp}(\alpha)$, then $X = e^Y \sim \text{Pareto}(\alpha)$.

Normal distribution

- Let $X_1, \dots, X_n \stackrel{\text{i.i.d.}}{\sim} \mathcal{N}(\mu, \sigma^2)$.
Then $X_1 + \dots + X_n \sim \mathcal{N}(n\mu, n\sigma^2)$.

Probability transformations

- If two RVs $X_1, X_2 \sim F_X$, then the CDF of the RV $Y = \max\{X_1, X_2\}$ is $F_Y(y) = (F_X(y))^2$.

intentionally left blank

Notations

Unless otherwise specified, the following notations were used:

$\mathbf{1}_d$	identity vector ($\in \mathbb{R}^d$)	\mathcal{U}	uniform distribution
\mathbb{I}_d	identity matrix ($\in \mathbb{R}^{d \times d}$)	\mathcal{N}	stand. normal distr.
$\mathbb{I}_{\{A\}}$	indicator function	ϕ	standard normal PDF
		Φ	standard normal CDF

Abbreviations

a.s.	almost surely	MDA	maximum domain of attraction
CDF	cumulative distribution function	MLE	maximum likelihood estimator
CF	characteristic function	PDF	probability density function
c.m.	completely monotone	QRM	quantitative risk management
EDF	empirical density function (CDF)	RV	random variable
GEV	generalized extreme value	s.t.	such that
iff	if and only if	w/	with
i.i.d.	independent and identically distributed	w.l.o.g.	without loss of generality
IOT	in order to	w.r.t.	with respect to

Disclaimer

- This summary is work in progress, i.e. neither completeness nor correctness of the content are guaranteed by the author.
- This summary may be extended or modified at the discretion of the readers.
- *Source:* Lecture Quantitative Risk Management, Spring Semester 2016, ETHZ (lecture notes, script, exercises and literature). Copyright of the content is with the lecturers.
- The layout of this summary is based on the summaries of several courses of the BSc ETH ME from Jonas LIECHTI.